

RISC-V Bitmanip Extension

Document Version draft

Editor: Clifford Wolf
Symbiotic GmbH
`clifford@symbioticeda.com`
September 1, 2019

Contributors to all versions of the spec in alphabetical order (please contact editors to suggest corrections): Jacob Bachmeyer, Allen Baum, Ari Ben, Alex Bradbury, Steven Braeger, Rogier Brussee, Michael Clark, Ken Dockser, Paul Donahue, Dennis Ferguson, Fabian Giesen, John Hauser, Robert Henry, Bruce Houtt, Po-wei Huang, Ben Marshall, Rex McCrary, Lee Moore, Jiří Moravec, Samuel Neves, Markus Oberhumer, Christopher Olson, Nils Pipenbrinck, Joseph Rahmeh, Xue Saw, Tommy Thorn, Avishai Tvila, Andrew Waterman, Thomas Wicki, and Clifford Wolf.

This document is released under a Creative Commons Attribution 4.0 International License.

Contents

1	Introduction	1
1.1	ISA Extension Proposal Design Criteria	1
1.2	B Extension Adoption Strategy	2
1.3	Next steps	2
2	RISC-V Bitmanip Extension	3
2.1	Basic bit manipulation instructions	4
2.1.1	Count Leading/Trailing Zeros (clz , ctz)	4
2.1.2	Count Bits Set (pcnt)	5
2.1.3	Logic-with-negate (andn , orn , xnor)	5
2.1.4	Pack two XLEN/2 words in one register (pack)	6
2.1.5	Min/max instructions (min , max , minu , maxu)	7
2.1.6	Single-bit instructions (sbset , sbclr , sbinv , sbext)	8
2.1.7	Shift Ones (Left/Right) (slo , sloi , sro , sroi)	9
2.2	Bit permutation instructions	10
2.2.1	Rotate (Left/Right) (rol , ror , rori)	10
2.2.2	Generalized Reverse (grev , grevi , rev)	11
2.2.3	Generalized Shuffle (shfl , unshfl , shfli , unshfli , zip , unzip)	15
2.3	Generalized OR-Combine (gorc , gorci)	22
2.4	Bit-Field Place (bfp)	24
2.5	Bit Extract/Deposit (bext , bdep)	25

2.6	Carry-Less Multiply (<code>clmul</code> , <code>clmulh</code> , <code>clmulr</code>)	26
2.7	CRC Instructions (<code>crc32.[bhwd]</code> , <code>crc32c.[bhwd]</code>)	28
2.8	Bit-Matrix Instructions (<code>bmatxor</code> , <code>bmator</code> , <code>bmatflip</code> , RV64 only)	29
2.9	Ternary Bit-Manipulation Instructions	31
2.9.1	Conditional Mix (<code>cmix</code>)	31
2.9.2	Conditional Move (<code>cmov</code>)	32
2.9.3	Funnel Shift (<code>fsl</code> , <code>fsr</code> , <code>fsri</code>)	33
2.10	Unsigned address calculation instructions	34
2.10.1	Add/sub with postfix zero-extend (<code>addwu</code> , <code>subwu</code> , <code>addiwu</code>)	34
2.10.2	Add/sub/shift with prefix zero-extend (<code>addu.w</code> , <code>subu.w</code> , <code>slliu.w</code>)	35
2.11	Opcode Encodings	36
2.12	Future compressed instructions	41
2.13	Micro architectural considerations and macro-op fusion for bit-manipulation	41
2.13.1	Fast MUL, MULH, MULHSU, MULHU	41
2.13.2	Fused load-immediate sequences	41
2.13.3	Fused <code>*-bfp</code> sequences	43
2.13.4	Fused <code>*-not</code> sequences	43
2.13.5	Fused <code>*-srli</code> and <code>*-srai</code> sequences	43
2.13.6	Fused sequences for logic operations	44
2.13.7	Fused ternary ALU sequences	45
2.13.8	Pseudo-ops for fused sequences	45
2.14	C intrinsics via <code><rvintrin.h></code>	45
3	Reference Implementations	49
3.1	Verilog reference implementations	49
3.2	Fast C reference implementations	51
4	Evaluation	55
4.1	Basic Bitmanipulation	55

4.1.1	Bitfield extract	55
4.1.2	Parity check	55
4.1.3	Rank and select	56
4.1.4	Packing bytes in word	56
4.1.5	Counting trailing non-zero bytes	56
4.1.6	Finding bytes of certain values	57
4.1.7	Fill right of most significant set bit	57
4.1.8	Round to next power of two	58
4.2	Funnel shifts	59
4.2.1	Bigint shift	59
4.2.2	Parsing bit-streams	60
4.2.3	Fixed-point multiply	61
4.3	Arbitrary bit permutations	62
4.3.1	Using butterfly operations	62
4.3.2	Using omega-flip networks	63
4.3.3	Using baseline networks	63
4.3.4	Using sheep-and-goats	64
4.3.5	Using bit-matrix multiply	64
4.4	Mirroring and rotating bitboards	65
4.4.1	Mirroring bitboards	65
4.4.2	Rotating bitboards	66
4.4.3	Explanation	66
4.4.4	Rotating Bitcubes	67
4.5	Inverting Xorshift RNGs	67
4.6	Cyclic redundancy checks (CRC)	69
4.7	Decoding RISC-V Immediates	71

Bibliography**75**

Chapter 1

Introduction

This is the RISC-V Bitmanip Extension draft spec.

1.1 ISA Extension Proposal Design Criteria

Any proposed changes to the ISA should be evaluated according to the following criteria.

- **Architecture Consistency:** Decisions must be consistent with RISC-V philosophy. ISA changes should deviate as little as possible from existing RISC-V standards (such as instruction encodings), and should not re-implement features that are already found in the base specification or other extensions.
- **Threshold Metric:** The proposal should provide *significant* savings in terms of clocks or instructions. As a heuristic, any proposal should replace at least three instructions. An instruction that only replaces two may be considered, but only if the frequency of use is very high and/or the implementation very cheap.
- **Data-Driven Value:** Usage in real world applications, and corresponding benchmarks showing a performance increase, will contribute to the score of a proposal. A proposal will not be accepted on the merits of its *theoretical* value alone, unless it is used in the real world.
- **Hardware Simplicity:** Though instructions saved is the primary benefit, proposals that dramatically increase the hardware complexity and area, or are difficult to implement, should be penalized and given extra scrutiny. The final proposals should only be made if a test implementation can be produced.
- **Compiler Support:** ISA changes that can be natively detected by the compiler, or are already used as intrinsics, will score higher than instructions which do not fit that criteria.

1.2 B Extension Adoption Strategy

The overall goal of this extension is pervasive adoption by minimizing potential barriers and ensuring the instructions can be mapped to the largest number of ops, either direct or pseudo, that are supported by the most popular processors and compilers. By adding generic instructions and taking advantage of the RISC-V base instructions that already operate on bits, the minimal set of instructions need to be added while at the same time enabling a rich of operations.

The instructions cover the four major categories of bit manipulation: Count, Extract, Insert, Swap. The spec supports RV32, RV64, and RV128. “Clever” obscure and/or overly specific instructions are avoided in favor of more straightforward, fast, generic ones. Coordination with other emerging RISC-V ISA extensions groups is required to ensure our instruction sets are architecturally consistent.

1.3 Next steps

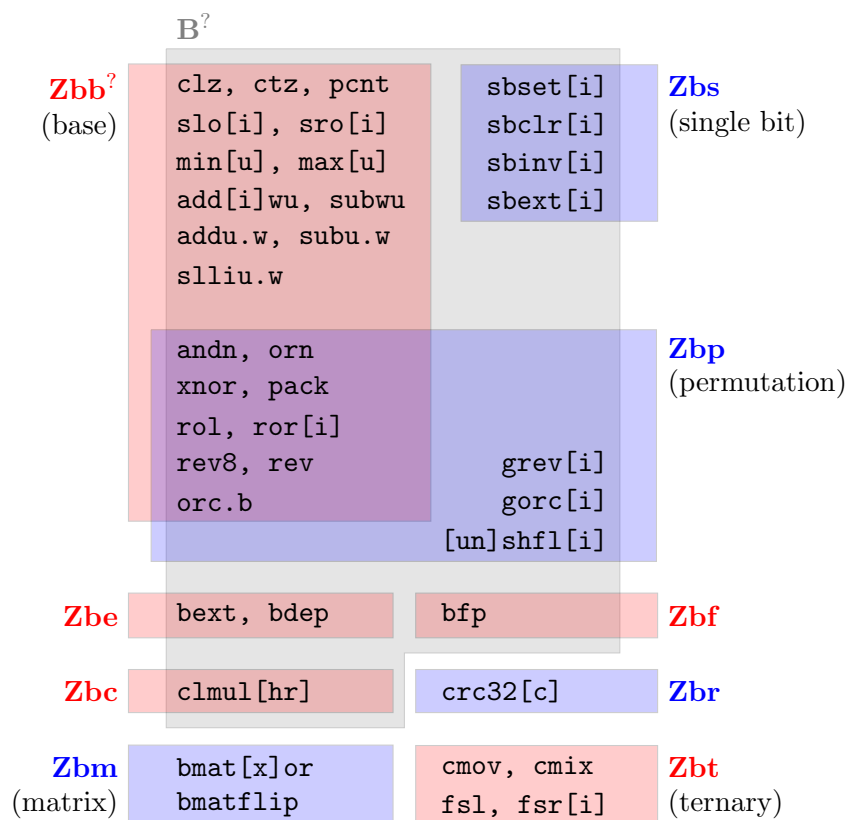
- Assign concrete instruction encodings so that we can start implementing the extension in processor cores and compilers.
- Add support for this extension to processor cores and compilers so we can run quantitative evaluations on the instructions.
- Create assembler snippets for common operations that do not map 1:1 to any instruction in this spec, but can be implemented easily using clever combinations of the instructions. Add support for those snippets to compilers.

Chapter 2

RISC-V Bitmanip Extension

In the proposals provided in this chapter, the C code examples are for illustration purposes only. They are not optimal implementations, but are intended to specify the desired functionality.

The final standard will likely define a range of Z-extensions for different bit manipulation instructions, with the “B” extension itself being a mix of instructions from those Z-extensions. It is unclear as of yet what this will look like exactly, but it will probably look something like this:



The main open questions of course relate to what should and shouldn't be included in “B”, and

what should or shouldn't be included in "Zbb". These decisions will be informed in big part by evaluations of the cost and added value for the individual instructions.

The main open questions are:

- Should `clmul[hr]` be included in "B", or `crc32.[bhwd]/crc32c.[bhwd]`, or neither, or both?
- Should "Zbe" be included in "B"? Should "Zbm be included in "B"?
- Which "Zbp" pseudo-ops should be included in "B"? Which in "Zbb"? Should "Zbp" be included in "B" as a whole?

For the purpose of tool-chain development "B" is currently everything.

For extensions that only implement certain pseudo-instructions (such as "Zbb" implements `rev8` and `rev`, which are pseudo-instructions for `grevi rd, rs1, -8` and `grevi rd, rs1, -1` respectively, the same binary encoding is used for those instructions as are used on a core with full support for the `grev[i]` instruction.

2.1 Basic bit manipulation instructions

2.1.1 Count Leading/Trailing Zeros (`clz`, `ctz`)

RISC-V Bitmanip ISA
RV32, RV64: <code>clz rd, rs</code> <code>ctz rd, rs</code> RV64 only: <code>clzw rd, rs</code> <code>ctzw rd, rs</code>

The `clz` operation counts the number of 0 bits at the MSB end of the argument. That is, the number of 0 bits before the first 1 bit counting from the most significant bit. If the input is 0, the output is XLEN. If the input is -1, the output is 0.

The `ctz` operation counts the number of 0 bits at the LSB end of the argument. If the input is 0, the output is XLEN. If the input is -1, the output is 0.

```
uint_xlen_t clz(uint_xlen_t rs1)
{
    for (int count = 0; count < XLEN; count++)
        if ((rs1 << count) >> (XLEN - 1))
            return count;
    return XLEN;
}
```

```

uint_xlen_t ctz(uint_xlen_t rs1)
{
    for (int count = 0; count < XLEN; count++)
        if ((rs1 >> count) & 1)
            return count;
    return XLEN;
}

```

The expression `XLEN-1-clz(x)` evaluates to the index of the most significant set bit, also known as integer base-2 logarithm, or -1 if `x` is zero.

2.1.2 Count Bits Set (pcnt)

	RISC-V Bitmanip ISA
RV32, RV64:	
pcnt rd, rs	
RV64 only:	
pcntw rd, rs	

This instruction counts the number of 1 bits in a register. This operations is known as population count, popcount, sideways sum, bit summation, or Hamming weight. [22, 20]

```

uint_xlen_t pcnt(uint_xlen_t rs1)
{
    int count = 0;
    for (int index = 0; index < XLEN; index++)
        count += (rs1 >> index) & 1;
    return count;
}

```

2.1.3 Logic-with-negate (andn, orn, xnor)

	RISC-V Bitmanip ISA
RV32, RV64:	
andn rd, rs1, rs2	
orn rd, rs1, rs2	
xnor rd, rs1, rs2	

This instructions implement AND, OR, and XOR with the 2nd arument inverted.

```

uint_xlen_t andn(uint_xlen_t rs1, uint_xlen_t rs2)
{
    return rs1 & ~rs2;
}

```

```

uint_xlen_t orn(uint_xlen_t rs1, uint_xlen_t rs2)
{
    return rs1 | ~rs2;
}

uint_xlen_t xnor(uint_xlen_t rs1, uint_xlen_t rs2)
{
    return rs1 ^ ~rs2;
}

```

This can use the existing inverter on rs2 in the ALU that's already there to implement subtract.

Among other things, those instructions allow implementing the “trailing bit manipulation” code patterns in two instructions each. For example, $(x - 1) \& \sim x$ produces a mask from trailing zero bits in x .

2.1.4 Pack two XLEN/2 words in one register (pack)

RISC-V Bitmanip ISA	
RV32, RV64:	<code>pack rd, rs1, rs2</code>
RV64 only:	<code>packw rd, rs1, rs2</code>

This instruction packs the XLEN/2-bit lower halves of rs1 and rs2 into rd, with rs1 in the lower half and rs2 in the upper half.

```

uint_xlen_t pack(uint_xlen_t rs1, uint_xlen_t rs2)
{
    uint_xlen_t lower = (rs1 << XLEN/2) >> XLEN/2;
    uint_xlen_t upper = rs2 << XLEN/2;
    return upper | lower;
}

```

Applications include XLEN/2-bit funnel shifts, zero-extend XLEN/2 bit values, duplicate the lower XLEN/2 bits (e.g. for mask creation), and loading unsigned 32 constants on RV64.

```

; Load 0xffff0000ffff0000 on RV64
lui rd, 0xffff0
pack rd, rd, rd

; Same as FSLW on RV64
pack rd, rs1, rs3
rol rd, rd, rs2
addiw rd, rd, 0

; Clear the upper half of rd
pack rd, rd, zero

```

Paired with `shfli/unshfli` and the other bit permutation instructions, `pack` can interleave arbitrary power-of-two chunks of `rs1` and `rs2`. For example, interleaving the bytes in the lower halves of `rs1` and `rs2`:

```
pack rd, rs1, rs2
zip8 rd, rd
```

`pack` is most commonly used to zero-extend words $< \text{XLEN}$. For this purpose we define the following assembler pseudo-ops:

RV32:

```
zext.b rd, rs  ->  andi  rd, rs, 255
zext.h rd, rs  ->  pack  rd, rs, zero
```

RV64:

```
zext.b rd, rs  ->  andi  rd, rs, 255
zext.h rd, rs  ->  packw rd, rs, zero
zext.w rd, rs  ->  pack  rd, rs, zero
```

RV128:.

```
zext.b rd, rs  ->  andi  rd, rs, 255
zext.h rd, rs  ->  packw rd, rs, zero
zext.w rd, rs  ->  packd rd, rs, zero
zext.d rd, rs  ->  pack  rd, rs, zero
```

2.1.5 Min/max instructions (`min`, `max`, `minu`, `maxu`)

RISC-V Bitmanip ISA

RV32, RV64:

```
min  rd, rs1, rs2
max  rd, rs1, rs2
minu rd, rs1, rs2
maxu rd, rs1, rs2
```

We define 4 R-type instructions `min`, `max`, `minu`, `maxu` with the following semantics:

```
uint_xlen_t min(uint_xlen_t rs1, uint_xlen_t rs2)
{
    return (int_xlen_t)rs1 < (int_xlen_t)rs2 ? rs1 : rs2;
}

uint_xlen_t max(uint_xlen_t rs1, uint_xlen_t rs2)
{
    return (int_xlen_t)rs1 > (int_xlen_t)rs2 ? rs1 : rs2;
}

uint_xlen_t minu(uint_xlen_t rs1, uint_xlen_t rs2)
{
    return rs1 < rs2 ? rs1 : rs2;
}
```

```
uint_xlen_t maxu(uint_xlen_t rs1, uint_xlen_t rs2)
{
    return rs1 > rs2 ? rs1 : rs2;
}
```

Code that performs saturated arithmetic on a word size $< \text{XLEN}$ needs to perform min/max operations frequently. A simple way of performing those operations without branching can benefit those programs.

SAT solvers spend a lot of time calculating the absolute value of a signed integer due to the way CNF literals are commonly encoded [10]. With `max` (or `minu`) this is a two-instruction operation:

```
neg a1, a0
max a0, a0, a1
```

2.1.6 Single-bit instructions (`sbset`, `sbclr`, `sbinv`, `sbext`)

RISC-V Bitmanip ISA

RV32, RV64:

```
sbset  rd, rs1, rs2
sbclr  rd, rs1, rs2
sbinv  rd, rs1, rs2
sbext  rd, rs1, rs2
sbseti rd, rs1, imm
sbclri rd, rs1, imm
sbinvi rd, rs1, imm
sbexti rd, rs1, imm
```

RV64:

```
sbsetw  rd, rs1, rs2
sbclrw  rd, rs1, rs2
sbinvw  rd, rs1, rs2
sbextw  rd, rs1, rs2
sbsetiw rd, rs1, imm
sbclriw rd, rs1, imm
sbinviw rd, rs1, imm
```

We define 4 single-bit instructions `sbset` (set), `sbclr` (clear), `sbinv` (invert), and `sbext` (extract), and their immediate-variants, with the following semantics:

```
uint_xlen_t sbset(uint_xlen_t rs1, uint_xlen_t rs2)
{
    int shamt = rs2 & (XLEN - 1);
    return rs1 | (uint_xlen_t(1) << shamt);
}
```

```

uint_xlen_t sbclr(uint_xlen_t rs1, uint_xlen_t rs2)
{
    int shamt = rs2 & (XLEN - 1);
    return rs1 & ~(uint_xlen_t(1) << shamt);
}

uint_xlen_t sbinv(uint_xlen_t rs1, uint_xlen_t rs2)
{
    int shamt = rs2 & (XLEN - 1);
    return rs1 ^ (uint_xlen_t(1) << shamt);
}

uint_xlen_t sbext(uint_xlen_t rs1, uint_xlen_t rs2)
{
    int shamt = rs2 & (XLEN - 1);
    return 1 & (rs1 >> shamt);
}

```

2.1.7 Shift Ones (Left/Right) (slo, sloi, sro, sroi)

RISC-V Bitmanip ISA

```

RV32, RV64:
    slo rd, rs1, rs2
    sro rd, rs1, rs2
    sloi rd, rs1, imm
    sroi rd, rs1, imm

RV64 only:
    slow rd, rs1, rs2
    srow rd, rs1, rs2
    sloiw rd, rs1, imm
    sroiw rd, rs1, imm

```

These instructions are similar to shift-logical operations from the base spec, except instead of shifting in zeros, they shift in ones.

```

uint_xlen_t slo(uint_xlen_t rs1, uint_xlen_t rs2)
{
    int shamt = rs2 & (XLEN - 1);
    return ~(~rs1 << shamt);
}

uint_xlen_t sro(uint_xlen_t rs1, uint_xlen_t rs2)
{
    int shamt = rs2 & (XLEN - 1);
    return ~(~rs1 >> shamt);
}

```

ISAs with flag registers often have a "Shift in Carry" or "Rotate through Carry" instruction.

Arguably a "Shift Ones" is an equivalent on an ISA like RISC-V that avoids such flag registers.

The main application for the Shift Ones instruction is mask generation.

When implementing this circuit, the only change in the ALU over a standard logical shift is that the value shifted in is not zero, but is a 1-bit register value that has been forwarded from the high bit of the instruction decode. This creates the desired behavior on both logical zero-shifts and logical ones-shifts.

2.2 Bit permutation instructions

2.2.1 Rotate (Left/Right) (rol, ror, rori)

RISC-V Bitmanip ISA

RV32, RV64:

```
ror rd, rs1, rs2
rol rd, rs1, rs2
rori rd, rs1, imm
```

RV64 only:

```
rorw rd, rs1, rs2
rolw rd, rs1, rs2
roriw rd, rs1, imm
```

These instructions are similar to shift-logical operations from the base spec, except they shift in the values from the opposite side of the register, in order. This is also called 'circular shift'.

```
uint_xlen_t rol(uint_xlen_t rs1, uint_xlen_t rs2)
{
    int shamt = rs2 & (XLEN - 1);
    return (rs1 << shamt) | (rs1 >> ((XLEN - shamt) & (XLEN - 1)));
}

uint_xlen_t ror(uint_xlen_t rs1, uint_xlen_t rs2)
{
    int shamt = rs2 & (XLEN - 1);
    return (rs1 >> shamt) | (rs1 << ((XLEN - shamt) & (XLEN - 1)));
}
```

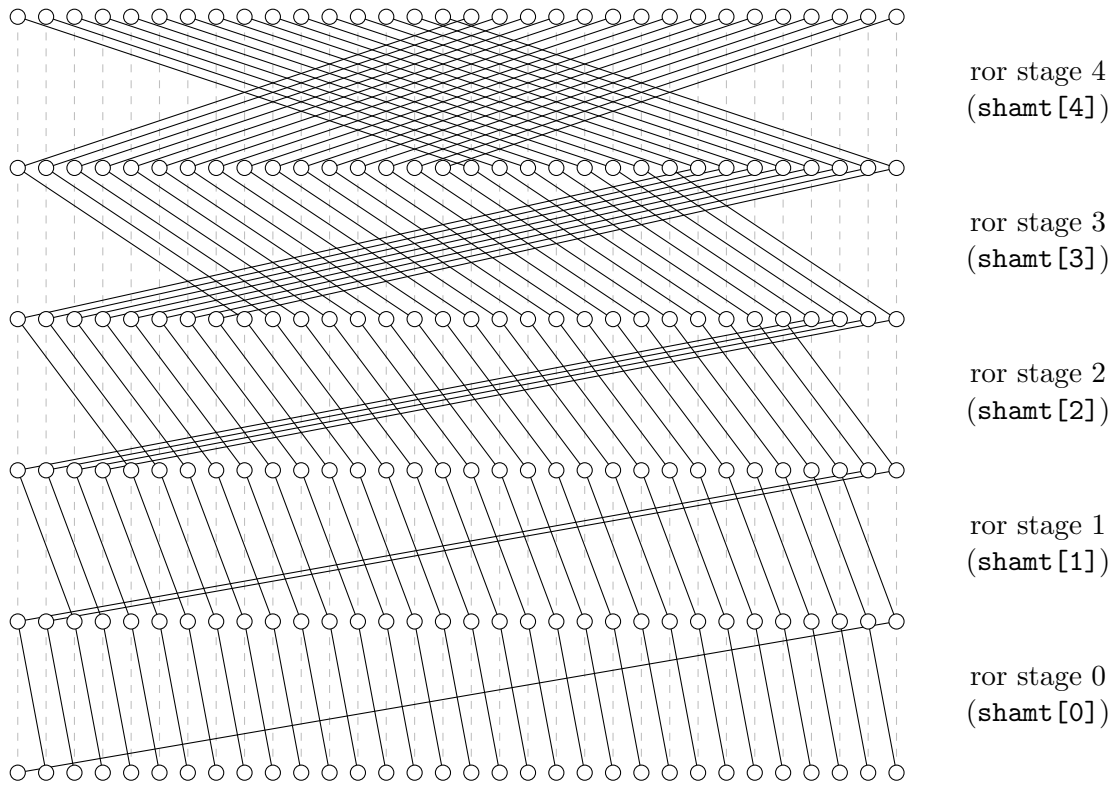



Figure 2.1: ror permutation network

2.2.2 Generalized Reverse (grev, grevi, rev)

RISC-V Bitmanip ISA	
RV32, RV64:	
<code>grev rd, rs1, rs2</code>	
<code>grevi rd, rs1, imm</code>	
RV64 only:	
<code>grevw rd, rs1, rs2</code>	
<code>greviw rd, rs1, imm</code>	

This instruction provides a single hardware instruction that can implement all of byte-order swap, bitwise reversal, short-order-swap, word-order-swap (RV64), nibble-order swap, bitwise reversal in a byte, etc, all from a single hardware instruction.

The Generalized Reverse (GREV) operation iteratively checks each bit i in the 2nd argument from $i = 0$ to $\log_2(\text{XLEN}) - 1$, and if the corresponding bit is set, swaps each adjacent pair of 2^i bits.




Figure 2.2: grev permutation network

```

uint32_t grev32(uint32_t rs1, uint32_t rs2)
{
    uint32_t x = rs1;
    int shamt = rs2 & 31;
    if (shamt & 1) x = ((x & 0x55555555) << 1) | ((x & 0xAAAAAAAA) >> 1);
    if (shamt & 2) x = ((x & 0x33333333) << 2) | ((x & 0xCCCCCCCC) >> 2);
    if (shamt & 4) x = ((x & 0x0F0F0F0F) << 4) | ((x & 0xF0F0F0F0) >> 4);
    if (shamt & 8) x = ((x & 0x00FF00FF) << 8) | ((x & 0xFF00FF00) >> 8);
    if (shamt & 16) x = ((x & 0x0000FFFF) << 16) | ((x & 0xFFFF0000) >> 16);
    return x;
}

```

```

uint64_t grev64(uint64_t rs1, uint64_t rs2)
{
    uint64_t x = rs1;
    int shamt = rs2 & 63;
    if (shamt & 1) x = ((x & 0x5555555555555555LL) << 1) |
                      ((x & 0xAAAAAAAAAAAAAAAAALL) >> 1);
    if (shamt & 2) x = ((x & 0x3333333333333333LL) << 2) |
                      ((x & 0xCCCCCCCCCCCCCCCCLL) >> 2);
    if (shamt & 4) x = ((x & 0x0F0F0F0F0F0F0F0FLL) << 4) |
                      ((x & 0xF0F0F0F0F0F0F0F0LL) >> 4);
    if (shamt & 8) x = ((x & 0x00FF00FF00FF00FFLL) << 8) |
                      ((x & 0xFF00FF00FF00FF00LL) >> 8);
    if (shamt & 16) x = ((x & 0x0000FFFF0000FFFFLL) << 16) |
                      ((x & 0xFFFF0000FFFF0000LL) >> 16);
    if (shamt & 32) x = ((x & 0x00000000FFFFFFFFLL) << 32) |
                      ((x & 0xFFFFFFFF00000000LL) >> 32);
    return x;
}

```

The above pattern should be intuitive to understand in order to extend this definition in an obvious manner for RV128.

The **grev** operation can easily be implemented using a permutation network with $\log_2(\text{XLEN})$ stages. Figure 2.1 shows the permutation network for **ror** for reference. Figure 2.2 shows the permutation network for **grev**.

Pseudo-instructions are provided for the most common GREVI use-cases. Their names consist of a prefix and an optional suffix. Each prefix and suffix corresponds to a bit mask (see Table 2.2). The GREVI control word is obtained by AND-ing the two masks together.

In other words, the prefix controls the number of zero bits at the LSB end of the control word, and the suffix controls the number of zeros at the MSB end of the control word.

rev8 reverses the order of bytes in a word, thus performs endianness conversion. This is equivalent to the ARM **REV** instructions or **BSWAP** on x86. ARM also has instructions for swapping the bytes in 16-bit and 32-bit words, and reversing the bit order (see table 2.3).

RV32		RV64			
shamt	Instruction	shamt	Instruction	shamt	Instruction
0: 00000	—	0: 000000	—	32: 100000	rev32
1: 00001	rev.p	1: 000001	rev.p	33: 100001	—
2: 00010	rev2.n	2: 000010	rev2.n	34: 100010	—
3: 00011	rev.n	3: 000011	rev.n	35: 100011	—
4: 00100	rev4.b	4: 000100	rev4.b	36: 100100	—
5: 00101	—	5: 000101	—	37: 100101	—
6: 00110	rev2.b	6: 000110	rev2.b	38: 100110	—
7: 00111	rev.b	7: 000111	rev.b	39: 100111	—
8: 01000	rev8.h	8: 001000	rev8.h	40: 101000	—
9: 01001	—	9: 001001	—	41: 101001	—
10: 01010	—	10: 001010	—	42: 101010	—
11: 01011	—	11: 001011	—	43: 101011	—
12: 01100	rev4.h	12: 001100	rev4.h	44: 101100	—
13: 01101	—	13: 001101	—	45: 101101	—
14: 01110	rev2.h	14: 001110	rev2.h	46: 101110	—
15: 01111	rev.h	15: 001111	rev.h	47: 101111	—
16: 10000	rev16	16: 010000	rev16.w	48: 110000	rev16
17: 10001	—	17: 010001	—	49: 110001	—
18: 10010	—	18: 010010	—	50: 110010	—
19: 10011	—	19: 010011	—	51: 110011	—
20: 10100	—	20: 010100	—	52: 110100	—
21: 10101	—	21: 010101	—	53: 110101	—
22: 10110	—	22: 010110	—	54: 110110	—
23: 10111	—	23: 010111	—	55: 110111	—
24: 11000	rev8	24: 011000	rev8.w	56: 111000	rev8
25: 11001	—	25: 011001	—	57: 111001	—
26: 11010	—	26: 011010	—	58: 111010	—
27: 11011	—	27: 011011	—	59: 111011	—
28: 11100	rev4	28: 011100	rev4.w	60: 111100	rev4
29: 11101	—	29: 011101	—	61: 111101	—
30: 11110	rev2	30: 011110	rev2.w	62: 111110	rev2
31: 11111	rev	31: 011111	rev.w	63: 111111	rev

Table 2.1: Pseudo-instructions for **grevi** instruction

Prefix	Mask	Suffix	Mask
rev	111111	—	111111
rev2	111110	.w	011111 (w = word)
rev4	111100	.h	001111 (h = half word)
rev8	111000	.b	000111 (b = byte)
rev16	110000	.n	000011 (n = nibble)
rev32	100000	.p	000001 (p = pair)

Table 2.2: Naming scheme for **grevi** pseudo-instructions. The prefix and suffix masks are ANDed to compute the immediate argument.

RISC-V	ARM	X86
rev	RBIT	—
rev8.h	REV16	—
rev8.w	REV32	—
rev8	REV	BSWAP

Table 2.3: Comparison of bit/byte reversal instructions

2.2.3 Generalized Shuffle (shfl, unshfl, shfli, unshfli, zip, unzip)

RISC-V Bitmanip ISA	
RV32, RV64:	
shfl	rd, rs1, rs2
unshfl	rd, rs1, rs2
shfli	rd, rs1, imm
unshfli	rd, rs1, imm
RV64 only:	
shflw	rd, rs1, rs2
unshflw	rd, rs1, rs2

Shuffle is the third bit permutation instruction in the RISC-V Bitmanip extension, after rotary shift and generalized reverse. It implements a generalization of the operation commonly known as perfect outer shuffle and its inverse (shuffle/unshuffle), also known as zip/unzip or interlace/uninterlace.

Bit permutations can be understood as reversible functions on bit indices (i.e. 5 bit functions on RV32 and 6 bit functions on RV64).

Operation	Corresponding function on bit indices
Rotate shift	Addition modulo XLEN
Generalized reverse	XOR with bitmask
Generalized shuffle	Bitpermutation

A generalized (un)shuffle operation has $\log_2(\text{XLEN}) - 1$ control bits, one for each pair of neighbouring bits in a bit index. When the bit is set, generalized shuffle will swap the two index bits. The **shfl** operation performs this swaps in MSB-to-LSB order (performing a rotate left shift on contiguous regions of set control bits), and the **unshfl** operation performs the swaps in LSB-to-MSB order (performing a rotate right shift on contiguous regions of set control bits). Combining up to $\log_2(\text{XLEN})$ of those **shfl**/**unshfl** operations can implement any bitpermutation on the bit indices.

The most common type of shuffle/unshuffle operation is one on an immediate control value that only contains one contiguous region of set bits. We call those operations zip/unzip and provide pseudo-instructions for them. The naming scheme for those pseudo-instructions is similar to the naming scheme for the **grevi** pseudo-instructions (see Tables 2.1 and 2.2), except that the LSB bit of the masks in Table 2.2 is not used for zip/unzip.

Shuffle/unshuffle operations that only have individual bits set (not a contiguous region of two or more bits) are their own inverse.

shamt	inv	Bit index rotations	Pseudo-Instruction
0: 0000	0	no-op	—
0000	1	no-op	—
1: 0001	0	i[1] -> i[0]	zip.n, unzip.n
0001	1	<i>equivalent to 0001 0</i>	—
2: 0010	0	i[2] -> i[1]	zip2.b, unzip2.b
0010	1	<i>equivalent to 0010 0</i>	—
3: 0011	0	i[2] -> i[0]	zip.b
0011	1	i[2] <- i[0]	unzip.b
4: 0100	0	i[3] -> i[2]	zip4.h, unzip4.h
0100	1	<i>equivalent to 0100 0</i>	—
5: 0101	0	i[3] -> i[2], i[1] -> i[0]	—
0101	1	<i>equivalent to 0101 0</i>	—
6: 0110	0	i[3] -> i[1]	zip2.h
0110	1	i[3] <- i[1]	unzip2.h
7: 0111	0	i[3] -> i[0]	zip.h
0111	1	i[3] <- i[0]	unzip.h
8: 1000	0	i[4] -> i[3]	zip8, unzip8
1000	1	<i>equivalent to 1000 0</i>	—
9: 1001	0	i[4] -> i[3], i[1] -> i[0]	—
1001	1	<i>equivalent to 1001 0</i>	—
10: 1010	0	i[4] -> i[3], i[2] -> i[1]	—
1010	1	<i>equivalent to 1010 0</i>	—
11: 1011	0	i[4] -> i[3], i[2] -> i[0]	—
1011	1	i[4] <- i[3], i[2] <- i[0]	—
12: 1100	0	i[4] -> i[2]	zip4
1100	1	i[4] <- i[2]	unzip4
13: 1101	0	i[4] -> i[2], i[1] -> i[0]	—
1101	1	i[4] <- i[2], i[1] <- i[0]	—
14: 1110	0	i[4] -> i[1]	zip2
1110	1	i[4] <- i[1]	unzip2
15: 1111	0	i[4] -> i[0]	zip
1111	1	i[4] <- i[0]	unzip

Table 2.4: RV32 modes and pseudo-instructions for shfli/unshfli instruction

Like GREV and rotate shift, the (un)shuffle instruction can be implemented using a short sequence of elementary permutations, that are enabled or disabled by the shamt bits. But (un)shuffle has one stage fewer than GREV. Thus shfli+unshfli together require the same amount of encoding space as grevi.

```
uint32_t shuffle32_stage(uint32_t src, uint32_t maskL, uint32_t maskR, int N)
{
    uint32_t x = src & ~(maskL | maskR);
    x |= ((src << N) & maskL) | ((src >> N) & maskR);
    return x;
}
```

shamt	inv	Pseudo-Instruction	shamt	inv	Pseudo-Instruction
0: 00000	0	—	16: 10000	0	zip16, unzip16
00000	1	—	10000	1	—
1: 00001	0	zip.n, unzip.n	17: 10001	0	—
00001	1	—	10001	1	—
2: 00010	0	zip2.b, unzip2.b	18: 10010	0	—
00010	1	—	10010	1	—
3: 00011	0	zip.b	19: 10011	0	—
00011	1	unzip.b	10011	1	—
4: 00100	0	zip4.h, unzip4.h	20: 10100	0	—
00100	1	—	10100	1	—
5: 00101	0	—	21: 10101	0	—
00101	1	—	10101	1	—
6: 00110	0	zip2.h	22: 10110	0	—
00110	1	unzip2.h	10110	1	—
7: 00111	0	zip.h	23: 10111	0	—
00111	1	unzip.h	10111	1	—
8: 01000	0	zip8.w, unzip8.w	24: 11000	0	zip8
01000	1	—	11000	1	unzip8
9: 01001	0	—	25: 11001	0	—
01001	1	—	11001	1	—
10: 01010	0	—	26: 11010	0	—
01010	1	—	11010	1	—
11: 01011	0	—	27: 11011	0	—
01011	1	—	11011	1	—
12: 01100	0	zip4.w	28: 11100	0	zip4
01100	1	unzip4.w	11100	1	unzip4
13: 01101	0	—	29: 11101	0	—
01101	1	—	11101	1	—
14: 01110	0	zip2.w	30: 11110	0	zip2
01110	1	unzip2.w	11110	1	unzip2
15: 01111	0	zip.w	31: 11111	0	zip
01111	1	unzip.w	11111	1	unzip

Table 2.5: RV64 modes and pseudo-instructions for shfli/unshfli instruction

```

uint32_t shfl32(uint32_t rs1, uint32_t rs2)
{
    uint32_t x = rs1;
    int shamt = rs2 & 15;

    if (shamt & 8) x = shuffle32_stage(x, 0x00ff0000, 0x0000ff00, 8);
    if (shamt & 4) x = shuffle32_stage(x, 0x0f000f00, 0x00f000f0, 4);
    if (shamt & 2) x = shuffle32_stage(x, 0x30303030, 0x0c0c0c0c, 2);
    if (shamt & 1) x = shuffle32_stage(x, 0x44444444, 0x22222222, 1);

    return x;
}

```




Figure 2.3: (un)shuffle permutation network without “flip” stages

```

uint32_t unshfl32(uint32_t rs1, uint32_t rs2)
{
    uint32_t x = rs1;
    int shamt = rs2 & 15;

    if (shamt & 1) x = shuffle32_stage(x, 0x44444444, 0x22222222, 1);
    if (shamt & 2) x = shuffle32_stage(x, 0x30303030, 0x0c0c0c0c, 2);
    if (shamt & 4) x = shuffle32_stage(x, 0x0f000f00, 0x00f000f0, 4);
    if (shamt & 8) x = shuffle32_stage(x, 0x00ff0000, 0x0000ff00, 8);

    return x;
}

```

Or for RV64:

```

uint64_t shuffle64_stage(uint64_t src, uint64_t maskL, uint64_t maskR, int N)
{
    uint64_t x = src & ~(maskL | maskR);
    x |= ((src << N) & maskL) | ((src >> N) & maskR);
    return x;
}

```



```

uint64_t shfl64(uint64_t rs1, uint64_t rs2)
{
    uint64_t x = rs1;
    int shamt = rs2 & 31;

    if (shamt & 16) x = shuffle64_stage(x, 0x0000ffff00000000LL,
                                         0x00000000ffff0000LL, 16);
    if (shamt & 8) x = shuffle64_stage(x, 0x00ff000000ff0000LL,
                                       0x0000ff000000ff00LL, 8);
    if (shamt & 4) x = shuffle64_stage(x, 0x0f000f000f000f00LL,
                                       0x00f000f000f000f0LL, 4);
    if (shamt & 2) x = shuffle64_stage(x, 0x3030303030303030LL,
                                       0x0c0c0c0c0c0c0c0cLL, 2);
    if (shamt & 1) x = shuffle64_stage(x, 0x4444444444444444LL,
                                       0x2222222222222222LL, 1);

    return x;
}

uint64_t unshfl64(uint64_t rs1, uint64_t rs2)
{
    uint64_t x = rs1;
    int shamt = rs2 & 31;

    if (shamt & 1) x = shuffle64_stage(x, 0x4444444444444444LL,
                                         0x2222222222222222LL, 1);
    if (shamt & 2) x = shuffle64_stage(x, 0x3030303030303030LL,
                                       0x0c0c0c0c0c0c0c0cLL, 2);
    if (shamt & 4) x = shuffle64_stage(x, 0x0f000f000f000f00LL,
                                       0x00f000f000f000f0LL, 4);
    if (shamt & 8) x = shuffle64_stage(x, 0x00ff000000ff0000LL,
                                       0x0000ff000000ff00LL, 8);
    if (shamt & 16) x = shuffle64_stage(x, 0x0000ffff00000000LL,
                                       0x00000000ffff0000LL, 16);

    return x;
}

```

The above pattern should be intuitive to understand in order to extend this definition in an obvious manner for RV128.

Alternatively (un)shuffle can be implemented in a single network with one more stage than GREV, with the additional first and last stage executing a permutation that effectively reverses the order of the inner stages. However, since the inner stages only mux half of the bits in the word each, a hardware implementation using this additional “flip” stages might actually be more expensive than simply creating two networks.

```

uint32_t shuffle32_flip(uint32_t src)
{
    uint32_t x = src & 0x88224411;
    x |= ((src << 6) & 0x22001100) | ((src >> 6) & 0x00880044);
    x |= ((src << 9) & 0x00440000) | ((src >> 9) & 0x00002200);
    x |= ((src << 15) & 0x44110000) | ((src >> 15) & 0x00008822);
    x |= ((src << 21) & 0x11000000) | ((src >> 21) & 0x00000088);
    return x;
}

uint32_t unshfl32alt(uint32_t rs1, uint32_t rs2)
{
    uint32_t shfl_mode = 0;
    if (rs2 & 1) shfl_mode |= 8;
    if (rs2 & 2) shfl_mode |= 4;
    if (rs2 & 4) shfl_mode |= 2;
    if (rs2 & 8) shfl_mode |= 1;

    uint32_t x = rs1;
    x = shuffle32_flip(x);
    x = shfl32(x, shfl_mode);
    x = shuffle32_flip(x);

    return x;
}

```

Figure 2.4 shows the (un)shuffle permutation network with “flip” stages and Figure 2.3 shows the (un)shuffle permutation network without “flip” stages.

The `zip` instruction with the upper half of its input cleared performs the commonly needed “fan-out” operation. (Equivalent to `bdep` with a `0x55555555` mask.) The `zip` instruction applied twice fans out the bits in the lower quarter of the input word by a spacing of 4 bits.

For example, the following code calculates the bitwise prefix sum of the bits in the lower byte of a 32 bit word on RV32:

```

andi a0, a0, 0xff
zip a0, a0
zip a0, a0
slli a1, a0, 4
c.add a0, a1
slli a1, a0, 8
c.add a0, a1
slli a1, a0, 16
c.add a0, a1

```

The final prefix sum is stored in the 8 nibbles of the `a0` output word.

Similarly, the following code stores the indices of the set bits in the LSB nibbles of the output word (with the LSB bit having index 1), with the unused MSB nibbles in the output set to zero:




Figure 2.4: (un)shuffle permutation network with “flip” stages

```

andi a0, a0, 0xff
zip a0, a0
zip a0, a0
slli a1, a0, 1
or a0, a0, a1
slli a1, a0, 2
or a0, a0, a1
li a1, 0x87654321
and a1, a0, a1
bext a0, a1, a0

```

Other `zip` modes can be used to “fan-out” in blocks of 2, 4, 8, or 16 bit. `zip` can be combined with `grevi` to perform inner shuffles. For example on RV64:

```

li a0, 0x0000000012345678
zip4 t0, a0 ; <- 0x0102030405060708
rev4.b t1, t0 ; <- 0x1020304050607080
zip8 t2, a0 ; <- 0x0012003400560078
rev8.h t3, t2 ; <- 0x1200340056007800
zip16 t4, a0 ; <- 0x0000123400005678
rev16.w t5, t4 ; <- 0x1234000056780000

```

Another application for the `zip` instruction is generating Morton code [23].

The x86 `PUNPCK[LH] * MMX/SSE/AVX` instructions perform similar operations as `zip8` and `zip16`.

RV32		RV64			
shamt	Instruction	shamt	Instruction	shamt	Instruction
0: 00000	—	0: 000000	—	32: 100000	orc32
1: 00001	orc.p	1: 000001	orc.p	33: 100001	—
2: 00010	orc2.n	2: 000010	orc2.n	34: 100010	—
3: 00011	orc.n	3: 000011	orc.n	35: 100011	—
4: 00100	orc4.b	4: 000100	orc4.b	36: 100100	—
5: 00101	—	5: 000101	—	37: 100101	—
6: 00110	orc2.b	6: 000110	orc2.b	38: 100110	—
7: 00111	orc.b	7: 000111	orc.b	39: 100111	—
8: 01000	orc8.h	8: 001000	orc8.h	40: 101000	—
9: 01001	—	9: 001001	—	41: 101001	—
10: 01010	—	10: 001010	—	42: 101010	—
11: 01011	—	11: 001011	—	43: 101011	—
12: 01100	orc4.h	12: 001100	orc4.h	44: 101100	—
13: 01101	—	13: 001101	—	45: 101101	—
14: 01110	orc2.h	14: 001110	orc2.h	46: 101110	—
15: 01111	orc.h	15: 001111	orc.h	47: 101111	—
16: 10000	orc16	16: 010000	orc16.w	48: 110000	orc16
17: 10001	—	17: 010001	—	49: 110001	—
18: 10010	—	18: 010010	—	50: 110010	—
19: 10011	—	19: 010011	—	51: 110011	—
20: 10100	—	20: 010100	—	52: 110100	—
21: 10101	—	21: 010101	—	53: 110101	—
22: 10110	—	22: 010110	—	54: 110110	—
23: 10111	—	23: 010111	—	55: 110111	—
24: 11000	orc8	24: 011000	orc8.w	56: 111000	orc8
25: 11001	—	25: 011001	—	57: 111001	—
26: 11010	—	26: 011010	—	58: 111010	—
27: 11011	—	27: 011011	—	59: 111011	—
28: 11100	orc4	28: 011100	orc4.w	60: 111100	orc4
29: 11101	—	29: 011101	—	61: 111101	—
30: 11110	orc2	30: 011110	orc2.w	62: 111110	orc2
31: 11111	orc	31: 011111	orc.w	63: 111111	orc

Table 2.6: Pseudo-instructions for **gorci** instruction

2.3 Generalized OR-Combine (**gorc**, **gorci**)

RISC-V Bitmanip ISA	
RV32, RV64:	
gorc	rd, rs1, rs2
gorci	rd, rs1, imm
RV64 only:	
gorcw	rd, rs1, rs2
gorciw	rd, rs1, imm

The GORC operation is similar to GREV, except that instead of swapping pairs of bits, GORC ORs them together, and writes the new value in both positions.

```

uint32_t gorc32(uint32_t rs1, uint32_t rs2)
{
    uint32_t x = rs1;
    int shamt = rs2 & 31;
    if (shamt & 1) x |= ((x & 0x55555555) << 1) | ((x & 0xAAAAAAAA) >> 1);
    if (shamt & 2) x |= ((x & 0x33333333) << 2) | ((x & 0xCCCCCCCC) >> 2);
    if (shamt & 4) x |= ((x & 0x0F0F0F0F) << 4) | ((x & 0xF0F0F0F0) >> 4);
    if (shamt & 8) x |= ((x & 0x00FF00FF) << 8) | ((x & 0xFF00FF00) >> 8);
    if (shamt & 16) x |= ((x & 0x0000FFFF) << 16) | ((x & 0xFFFF0000) >> 16);
    return x;
}

uint64_t gorc64(uint64_t rs1, uint64_t rs2)
{
    uint64_t x = rs1;
    int shamt = rs2 & 63;
    if (shamt & 1) x |= ((x & 0x5555555555555555LL) << 1) |
        ((x & 0xAAAAAAAAAAAAAAAAALL) >> 1);
    if (shamt & 2) x |= ((x & 0x3333333333333333LL) << 2) |
        ((x & 0xCCCCCCCCCCCCCCCCLL) >> 2);
    if (shamt & 4) x |= ((x & 0x0F0F0F0F0F0F0F0FLL) << 4) |
        ((x & 0xF0F0F0F0F0F0F0F0LL) >> 4);
    if (shamt & 8) x |= ((x & 0x00FF00FF00FF00FFLL) << 8) |
        ((x & 0xFF00FF00FF00FF00LL) >> 8);
    if (shamt & 16) x |= ((x & 0x0000FFFF0000FFFFLL) << 16) |
        ((x & 0xFFFF0000FFFF0000LL) >> 16);
    if (shamt & 32) x |= ((x & 0x00000000FFFFFFFFLL) << 32) |
        ((x & 0xFFFFFFFF00000000LL) >> 32);
    return x;
}

```

GORC can be usefull for copying naturally aligned fields in a word, and testing such fields for being equal zero.

gorci pseudo-instructions follow the same naming scheme as **grevi** pseudo-instructions (see Tables 2.1 and 2.2), except the prefix **orc** is used instead of **rev**. See Table 2.6 for a full list of **gorci** pseudo-instructions.

An important use-case is **strlen()** and **strcpy()**, which can utilize **orc.b** for testing for zero bytes, and counting trailing non-zero bytes in a word.

2.4 Bit-Field Place (bfp)

RISC-V Bitmanip ISA

```
RV32, RV64:
    bfp rd, rs1, rs2

RV64 only:
    bfpw rd, rs1, rs2
```

The bit field place (**bfp**) instruction places up to 16 LSB bit from **rs2** into the value in **rs1**. The upper bits of **rs2** control the length of the bit field and target position.

```
uint_xlen_t bfp(uint_xlen_t rs1, uint_xlen_t rs2)
{
    int len = (rs2 >> 24) & 15;
    int off = (rs2 >> 16) & (XLEN-1);
    len = len ? len : 16;
    uint_xlen_t mask = rol(slo(0, len), off);
    uint_xlen_t data = rol(rs2, off);
    return (data & mask) | (rs1 & ~mask);
}
```

The layout of the control word in **rs2** is as follows. **LEN=0** encodes for **LEN=16**.

3			2																1																											
1	0	9	8	7	6	5	4	3	2	1	0	9	8	7	6	5	4	3	2	1	0	9	8	7	6	5	4	3	2	1	0															
-----																	-----																													
										LEN										OFF										DATA																

Placing bits from **a0** in **a1**, with results in **t0**:

```
addi t0, zero, {length[3:0], offset[7:0]}
pack t0, a0, t0
bfp t0, a1, t0
```

(On RV64 **packw** would be used as second instruction in that sequence.)

Placing up to 16 constant bits in any contiguous region:

```
lui t0, ...
addi t0, t0, ...
bfp t0, a1, t0
```

Note that either above sequence only modifies one register, which makes them fuse-able sequences.

2.5 Bit Extract/Deposit (bext, bdep)

RISC-V Bitmanip ISA

RV32, RV64:

```
bext rd, rs1, rs2
bdep rd, rs1, rs2
```

RV64 only:

```
bextw rd, rs1, rs2
bdepw rd, rs1, rs2
```

This instructions implement the generic bit extract and bit deposit functions. This operation is also referred to as bit gather/scatter, bit pack/unpack, parallel extract/deposit, compress/expand, or right_compress/right_expand.

bext collects LSB justified bits to rd from rs1 using extract mask in rs2.

bdep writes LSB justified bits from rs1 to rd using deposit mask in rs2.

```
uint_xlen_t bext(uint_xlen_t rs1, uint_xlen_t rs2)
{
    uint_xlen_t r = 0;
    for (int i = 0, j = 0; i < XLEN; i++)
        if ((rs2 >> i) & 1) {
            if ((rs1 >> i) & 1)
                r |= uint_xlen_t(1) << j;
            j++;
        }
    return r;
}

uint_xlen_t bdep(uint_xlen_t rs1, uint_xlen_t rs2)
{
    uint_xlen_t r = 0;
    for (int i = 0, j = 0; i < XLEN; i++)
        if ((rs2 >> i) & 1) {
            if ((rs1 >> j) & 1)
                r |= uint_xlen_t(1) << i;
            j++;
        }
    return r;
}
```

Implementations may choose to use smaller multi-cycle implementations of **bext** and **bdep**, or even emulate the instructions in software.

Even though multi-cycle **bext** and **bdep** often are not fast enough to outperform algorithms that use sequences of shifts and bit masks, dedicated instructions for those operations can still be of great advantage in cases where the mask argument is not constant.

For example, the following code efficiently calculates the index of the tenth set bit in `a0` using `bdep`:

```
li a1, 0x00000200
bdep a0, a1, a0
ctz a0, a0
```

For cases with a constant mask an optimizing compiler would decide when to use `bext` or `bdep` based on the optimization profile for the concrete processor it is optimizing for. This is similar to the decision whether to use `MUL` or `DIV` with a constant, or to perform the same operation using a longer sequence of much simpler operations.

The `bext` and `bdep` instructions are equivalent to the x86 BMI2 instructions `PEXT` and `PDEP`. But there is much older prior art. For example, the soviet BESM-6 mainframe computer, designed and built in the 1960s, had `APX/AUX` instructions with almost the same semantics. [1] (The BESM-6 `APX/AUX` instructions packed/unpacked at the MSB end instead of the LSB end. Otherwise it is the same instruction.)

Efficient hardware implementations of `bext` and `bdep` are described in [13] and demonstrated in [24].

2.6 Carry-Less Multiply (`clmul`, `clmulh`, `clmulr`)

RISC-V Bitmanip ISA

RV32, RV64:

```
clmul  rd, rs1, rs2
clmulh rd, rs1, rs2
clmulr rd, rs1, rs2
```

RV64 only:

```
clmulw  rd, rs1, rs2
clmulhw rd, rs1, rs2
clmulrw rd, rs1, rs2
```

Calculate the carry-less product [21] of the two arguments. `clmul` produces the lower half of the carry-less product and `clmulh` produces the upper half of the $2 \cdot \text{XLEN}$ carry-less product.

`clmulr` produces bits $2 \cdot \text{XLEN} - 2 : \text{XLEN} - 1$ of the $2 \cdot \text{XLEN}$ carry-less product. That means `clmulh` is equivalent to `clmulr` followed by a 1-bit right shift. (The MSB of a `clmulh` result is always zero.) Another equivalent definition of `clmulr` is that is `clmulr(a,b) := rev(clmul(rev(a), rev(b)))`. (The “r” in `clmulr` means reversed.)

Unlike `mulh[[s]u]`, we add a `*W` variant of `clmulh`. This is because we expect some code to use 32-bit `clmul` intrinsics, even on 64-bit architectures. For example in cases where data is processed in 32-bit chunks.


```

uint_xlen_t clmul(uint_xlen_t rs1, uint_xlen_t rs2)
{
    uint_xlen_t x = 0;
    for (int i = 0; i < XLEN; i++)
        if ((rs2 >> i) & 1)
            x ^= rs1 << i;
    return x;
}

uint_xlen_t clmulh(uint_xlen_t rs1, uint_xlen_t rs2)
{
    uint_xlen_t x = 0;
    for (int i = 1; i < XLEN; i++)
        if ((rs2 >> i) & 1)
            x ^= rs1 >> (XLEN-i);
    return x;
}

uint_xlen_t clmulr(uint_xlen_t rs1, uint_xlen_t rs2)
{
    uint_xlen_t x = 0;
    for (int i = 0; i < XLEN; i++)
        if ((rs2 >> i) & 1)
            x ^= rs1 >> (XLEN-i-1);
    return x;
}

```

The classic applications for `clmul` are CRC [11, 25] and GCM, but more applications exist, including the following examples.

There are obvious applications in hashing and pseudo random number generations. For example, it has been reported that hashes based on carry-less multiplications can outperform Google's CityHash [17].

`clmul` of a number with itself inserts zeroes between each input bit. This can be useful for generating Morton code [23].

`clmul` of a number with -1 calculates the prefix XOR operation. This can be useful for decoding gray codes.

Another application of XOR prefix sums calculated with `clmul` is branchless tracking of quoted strings in high-performance parsers. [16]

Carry-less multiply can also be used to implement Erasure code efficiently. [14]

SPARC introduced similar instructions (`XMULX`, `XMULXHI`) in SPARC T3 in 2010. [6]

TI C6000 introduced a similar instruction (`XORMPY`) in C64x+. [7]

2.7 CRC Instructions (`crc32.[bhwd]`, `crc32c.[bhwd]`)

RISC-V Bitmanip ISA

```
RV32, RV64:
    crc32.b rd, rs
    crc32.h rd, rs
    crc32.w rd, rs
    crc32c.b rd, rs
    crc32c.h rd, rs
    crc32c.w rd, rs

RV64 only:
    crc32.d rd, rs
    crc32c.d rd, rs
```

Unary CRC instructions that interpret the bits of `rs1` as a CRC32/CRC32C state and perform a polynomial reduction of that state shifted left by 8, 16, 32, or 64 bits.

The instructions return the new CRC32/CRC32C state.

The `crc32.w/crc32c.w` instructions are equivalent to executing `crc32.h/crc32c.h` twice, and `crc32.h/crc32c.h` instructions are equivalent to executing `crc32.b/crc32c.b` twice.

All 8 CRC instructions operate on bit-reflected data.

```
uint_xlen_t crc32(uint_xlen_t x, int nbits)
{
    for (int i = 0; i < nbits; i++)
        x = (x >> 1) ^ (0xEDB88320 & ~((x&1)-1));
    return x;
}

uint_xlen_t crc32c(uint_xlen_t x, int nbits)
{
    for (int i = 0; i < nbits; i++)
        x = (x >> 1) ^ (0x82F63B78 & ~((x&1)-1));
    return x;
}

uint_xlen_t crc32_b(uint_xlen_t rs1) { return crc32(rs1, 8); }
uint_xlen_t crc32_h(uint_xlen_t rs1) { return crc32(rs1, 16); }
uint_xlen_t crc32_w(uint_xlen_t rs1) { return crc32(rs1, 32); }

uint_xlen_t crc32c_b(uint_xlen_t rs1) { return crc32c(rs1, 8); }
uint_xlen_t crc32c_h(uint_xlen_t rs1) { return crc32c(rs1, 16); }
uint_xlen_t crc32c_w(uint_xlen_t rs1) { return crc32c(rs1, 32); }

#if XLEN > 32
uint_xlen_t crc32_d(uint_xlen_t rs1) { return crc32(rs1, 64); }
uint_xlen_t crc32c_d(uint_xlen_t rs1) { return crc32c(rs1, 64); }
#endif
```

Payload data must be XOR'ed into the LSB end of the state before executing the CRC instruction. The following code demonstrates the use of `crc32.b`:

```
uint32_t crc32_demo(const uint8_t *p, int len)
{
    uint32_t x = 0xffffffff;
    for (int i = 0; i < len; i++) {
        x = x ^ p[i];
        x = crc32_b(x);
    }
    return ~x;
}
```

In terms of binary polynomial arithmetic those instructions perform the operation

$$\text{rd}'(x) = (\text{rs1}'(x) \cdot x^N) \bmod \{1, P'\}(x),$$

with $N \in \{8, 16, 32, 64\}$, $P = 0\text{xEDB8_8320}$ for CRC32 and $P = 0\text{x82F6_3B78}$ for CRC32C, a' denoting the XLEN bit reversal of a , and $\{a, b\}$ denoting bit concatenation. Note that for example for CRC32 $\{1, P'\} = 0\text{x1_04C1_1DB7}$ on RV32 and $\{1, P'\} = 0\text{x1_04C1_1DB7_0000_0000}$ on RV64.

These dedicated CRC instructions are meant for RISC-V implementations without fast multiplier and therefore without fast `clmul[h]`. For implementations with fast `clmul[h]` it is recommended to use the methods described in [11] and demonstrated in [25] that can process XLEN input bits using just one carry-less multiply for arbitrary CRC polynomials.

In applications where those methods are not applicable it is possible to emulate the dedicated CRC instructions using two carry-less multiplies that implement a Barrett reduction. The following example implements a replacement for `crc32.w` (RV32).

```
crc32_w:
    li t0, 0xF7011641
    li t1, 0xEDB88320
    clmul a0, a0, t0
    clmulr a0, a0, t1
    ret
```

2.8 Bit-Matrix Instructions (`bmatxor`, `bmator`, `bmatflip`, RV64 only)

RISC-V Bitmanip ISA

RV64 only:

```
bmator rd, rs1, rs2
bmatxor rd, rs1, rs2
bmatflip rd, rs
```

These are 64-bit-only instruction that are not available on RV32. On RV128 they ignore the upper

half of operands and sign extend the results.

This instructions interpret a 64-bit value as 8x8 binary matrix.

bmatxor performs a matrix-matrix multiply with boolean AND as multiply operator and boolean XOR as addition operator.

bmator performs a matrix-matrix multiply with boolean AND as multiply operator and boolean OR as addition operator.

bmatflip is a unary operator that transposes the source matrix. It is equivalent to `zip; zip; zip` on RV64.

```
uint64_t bmatflip(uint64_t rs1)
{
    uint64_t x = rs1;
    x = shfl64(x, 31);
    x = shfl64(x, 31);
    x = shfl64(x, 31);
    return x;
}

uint64_t bmatxor(uint64_t rs1, uint64_t rs2)
{
    // transpose of rs2
    uint64_t rs2t = bmatflip(rs2);

    uint8_t u[8]; // rows of rs1
    uint8_t v[8]; // cols of rs2

    for (int i = 0; i < 8; i++) {
        u[i] = rs1 >> (i*8);
        v[i] = rs2t >> (i*8);
    }

    uint64_t x = 0;
    for (int i = 0; i < 64; i++) {
        if (pcnt(u[i / 8] & v[i % 8]) & 1)
            x |= 1LL << i;
    }

    return x;
}
```

```

uint64_t bmatxor(uint64_t rs1, uint64_t rs2)
{
    // transpose of rs2
    uint64_t rs2t = bmatflip(rs2);

    uint8_t u[8]; // rows of rs1
    uint8_t v[8]; // cols of rs2

    for (int i = 0; i < 8; i++) {
        u[i] = rs1 >> (i*8);
        v[i] = rs2t >> (i*8);
    }

    uint64_t x = 0;
    for (int i = 0; i < 64; i++) {
        if ((u[i / 8] & v[i % 8]) != 0)
            x |= 1LL << i;
    }

    return x;
}

```

Among other things, `bmatxor`/`bmatxor` can be used to perform arbitrary permutations of bits within each byte (permutation matrix as 2nd operand) or perform arbitrary permutations of bytes within a 64-bit word (permutation matrix as 1st operand).

There are similar instructions in Cray XMT [5]. The Cray X1 architecture even has a full 64x64 bit matrix multiply unit [4].

The MMIX architecture has MOR and MXOR instructions with the same semantic. [15, p. 182f]

The x86 EVEX/VEX/SSE instruction GF2P8AFFINEQB is equivalent to `bmatxor`.

The `bmm.8` instruction proposed in [12] is also equivalent to `bmatxor`.

2.9 Ternary Bit-Manipulation Instructions

2.9.1 Conditional Mix (`cmix`)

RISC-V Bitmanip ISA

RV32, RV64:
`cmix rd, rs2, rs1, rs3`

(Note that the assembler syntax of `cmix` has the `rs2` argument first to make assembler code more readable. But the reference C code code below uses the “architecturally correct” argument order `rs1, rs2, rs3`.)

The `cmix rd, rs2, rs1, rs3` instruction selects bits from `rs1` and `rs3` based on the bits in the control word `rs2`.

```
uint_xlen_t cmix(uint_xlen_t rs1, uint_xlen_t rs2, uint_xlen_t rs3)
{
    return (rs1 & rs2) | (rs3 & ~rs2);
}
```

It replaces sequences like the following.

```
and rd, rs1, rs2
andn t0, rs3, rs2
or rd, rd, t0
```

Using `cmix` a single butterfly stage can be implemented in only two instructions. Thus, arbitrary bit-permutations can be implemented using only 18 instructions (32 bit) or 22 instructions (64 bits).

2.9.2 Conditional Move (cmov)

RV32, RV64: <code>cmov rd, rs2, rs1, rs3</code>	RISC-V Bitmanip ISA
--	---------------------

(Note that the assembler syntax of `cmov` has the `rs2` argument first to make assembler code more readable. But the reference C code below uses the “architecturally correct” argument order `rs1, rs2, rs3`.)

The `cmov rd, rs2, rs1, rs3` instruction selects `rs1` if the control word `rs2` is non-zero, and `rs3` if the control word is zero.

```
uint_xlen_t cmov(uint_xlen_t rs1, uint_xlen_t rs2, uint_xlen_t rs3)
{
    return rs2 ? rs1 : rs3;
}
```

The `cmov` instruction helps avoiding branches, which can lead to better performance, and helps with constant-time code as used in some cryptography applications.

2.9.3 Funnel Shift (fsl, fsr, fsri)

RISC-V Bitmanip ISA

RV32, RV64:

```
fsl  rd, rs1, rs3, rs2
fsr  rd, rs1, rs3, rs2
fsri rd, rs1, rs3, imm
```

RV64 only:

```
fslw rd, rs1, rs3, rs2
fsrw rd, rs1, rs3, rs2
fsriw rd, rs1, rs3, imm
```

(Note that the assembler syntax for funnel shifts has the **rs2** argument last to make assembler code more readable. But the reference C code below uses the “architecturally correct” argument order **rs1, rs2, rs3**.)

The **fsl rd, rs1, rs3, rs2** instruction creates a $2 \cdot \text{XLEN}$ word by concatenating **rs1** and **rs3** (with **rs1** in the MSB half), rotate-left-shifts that word by the amount indicated in the $\log_2(\text{XLEN}) + 1$ LSB bits in **rs2**, and then writes the MSB half of the result to **rd**.

The **fsr rd, rs1, rs3, rs2** instruction creates a $2 \cdot \text{XLEN}$ word by concatenating **rs1** and **rs3** (with **rs1** in the LSB half), rotate-right-shifts that word by the amount indicated in the $\log_2(\text{XLEN}) + 1$ LSB bits in **rs2**, and then writes the LSB half of the result to **rd**.

```
uint_xlen_t fsl(uint_xlen_t rs1, uint_xlen_t rs2, uint_xlen_t rs3)
{
    int shamt = rs2 & (2*XLEN - 1);
    uint_xlen_t A = rs1, B = rs3;
    if (shamt >= XLEN) {
        shamt -= XLEN;
        A = rs3;
        B = rs1;
    }
    return shamt ? (A << shamt) | (B >> (XLEN-shamt)) : A;
}

uint_xlen_t fsr(uint_xlen_t rs1, uint_xlen_t rs2, uint_xlen_t rs3)
{
    int shamt = rs2 & (2*XLEN - 1);
    uint_xlen_t A = rs1, B = rs3;
    if (shamt >= XLEN) {
        shamt -= XLEN;
        A = rs3;
        B = rs1;
    }
    return shamt ? (A >> shamt) | (B << (XLEN-shamt)) : A;
}
```

A shift unit capable of either `fsl` or `fsr` is capable of performing all the other shift functions, including the other funnel shift, with only minimal additional logic.

For any values of A, B, and C:

```
fsl(A, B, C) = fsr(A, -B, C)
```

And for any values x and $0 \leq \text{shamt} < \text{XLEN}$:

```
sll(x, shamt) == fsl(x, shamt, 0)
srl(x, shamt) == fsr(x, shamt, 0)
sra(x, shamt) == fsr(x, shamt, sext_x)
slo(x, shamt) == fsl(x, shamt, ~0)
sro(x, shamt) == fsr(x, shamt, ~0)
ror(x, shamt) == fsr(x, shamt, x)
rol(x, shamt) == fsl(x, shamt, x)
```

Furthermore an RV64 implementation of either `fsl` or `fsr` is capable of performing the *W versions of all shift operations with only a few gates of additional control logic.

On RV128 there is no `fsri` instruction. But there is `fsriw` and `fsrid`.

2.10 Unsigned address calculation instructions

Consider C code that's using unsigned 32-bit ints as array indices. For example:

```
char addiwu_demo(char *p, unsigned int i) {
    return p[i-1];
}

int slliuw_demo(int *p, unsigned int i, unsigned int j) {
    return p[i^j];
}
```

In both cases the expression within `p[...]` must overflow according to 32-bit arithmetic, then be zero-extended, and then this zero-extended result must be used in the address calculation.

The instructions below make sure that no explicit `zext.w` instruction is needed in those cases, to make sure there is no systematic performance penalty for code like shown above on RV64 compared to RV32.

2.10.1 Add/sub with postfix zero-extend (`addwu`, `subwu`, `addiwu`)

RV64:

```
addwu rd, rs1, rs2
subwu rd, rs1, rs2
addiwu rd, rs1, imm
```


These instructions are identical to `addw`, `subw`, `addiw`, except that bits XLEN-1:32 of the result are cleared after the addition. I.e. these instructions zero-extend instead of sign-extend the 32-bit result.

```
uint_xlen_t addwu(uint_xlen_t rs1, uint_xlen_t rs2)
{
    uint_xlen_t result = rs1 + rs2;
    return (uint32_t)result;
}

uint_xlen_t subwu(uint_xlen_t rs1, uint_xlen_t rs2)
{
    uint_xlen_t result = rs1 - rs2;
    return (uint32_t)result;
}
```

2.10.2 Add/sub/shift with prefix zero-extend (`addu.w`, `subu.w`, `slliu.w`)

RV64: <code>addu.w rd, rs1, rs2</code> <code>subu.w rd, rs1, rs2</code> <code>slliu.w rd, rs1, imm</code>
--

`slliu.w` is identical to `slli`, except that bits XLEN-1:32 of the `rs1` argument are cleared before the shift.

`addu.w` and `subu.w` are identical to `add` and `sub`, except that bits XLEN-1:32 of the `rs2` argument are cleared before the add/subtract.

```
uint_xlen_t slliuw(uint_xlen_t rs1, int imm)
{
    uint_xlen_t rs1u = (uint32_t)rs1;
    int shamt = imm & (XLEN - 1);
    return rs1u << shamt;
}

uint_xlen_t adduw(uint_xlen_t rs1, uint_xlen_t rs2)
{
    uint_xlen_t rs2u = (uint32_t)rs2;
    return rs1 + rs2u;
}

uint_xlen_t subuw(uint_xlen_t rs1, uint_xlen_t rs2)
{
    uint_xlen_t rs2u = (uint32_t)rs2;
    return rs1 - rs2u;
}
```

2.11 Opcode Encodings

This chapter contains proposed encodings for most of the instructions described in this document. **DO NOT IMPLEMENT THESE OPCODES YET.** We are trying to get official opcodes assigned and will update this chapter soon with the official opcodes.

The `andn`, `orn`, and `xnor` instruction are encoded the same way as `and`, `or`, and `xor`, but with `op[30]` set, mirroring the encoding scheme used for `add` and `sub`.

All shift instructions use `funct3=001` for left shifts and `funct3=101` for right shifts.

`op[26]=1` selects funnel shifts. For funnel shifts `op[30:29]` is part of the 3rd operand and therefore unused for encoding the operation. For all other shift operations `op[26]=0`.

`fsri` is also encoded with `op[26]=1`, leaving a 6 bit immediate. The 7th bit, that is necessary to perform a 128 bit funnel shift on RV64, can be emulated by swapping `rs1` and `rs3`.

There is no `shfliw` instruction. The `slliu.w` instruction occupies the encoding slot that would be occupied by `shfliw`.

On RV128 `op[26]` contains the MSB of the immediate for the shift instructions. Therefore there is no FSRI instruction on RV128. (But there is FSRIW/FSRID.)

	SLL	SRL	SRA	SLO	SRO	ROL	ROR	FSL	FSR
<code>op[30]</code>	0	0	1	0	0	1	1	-	-
<code>op[29]</code>	0	0	0	1	1	1	1	-	-
<code>op[26]</code>	0	0	0	0	0	0	0	1	1
<code>funct3</code>	001	101	101	001	101	001	101	001	101

Only an encoding for RORI exists, as ROLI can be implemented with RORI by negating the immediate. Unary functions are encoded in the spot that would correspond to ROLI, with the function encoded in the 5 LSB bits of the immediate.

The CRC instructions are encoded as unary instructions with `op[24]` set. The polynomial is selected via `op[23]`, with `op[23]=0` for CRC32 and `op[23]=1` for CRC32C. The width is selected with `op[22:20]`, using the same encoding as is used in `funct3` for load/store operations.

`cmix` and `cmov` are encoded using the two remaining ternary operator encodings in `funct3=001` and `funct3=101`. (There are two ternary operator encodings per minor opcode using the `op[26]=1` scheme for marking ternary OPs.)

The single-bit instructions are also encoded within the shift opcodes, with `op[27]` set, and using `op[30]` and `op[29]` to select the operation:

	SBCLR	SBSET	SBINV	SBEXT	GORC	GREV
<code>op[30]</code>	1	0	1	1	0	1
<code>op[29]</code>	0	1	1	0	1	1
<code>op[27]</code>	1	1	1	1	1	1
<code>funct3</code>	001	001	001	101	101	101

There is no `sbextiw` instruction as it can be emulated trivially using `sbexti`. However, there is

sbsetiw, **sbclriw**, and **sbinviw** as changing bit 31 would change the sign extend. There are non-immediate **W* instructions of all single-bit instructions, including **sbextw**, because the number of used bits in *rs2* is different in **sbext** and **sbextw**.

GORC and GREV are encoded in the two remaining slots in the single-bit instruction encoding space.

The remaining instructions are encoded within **funct7**=0000100 and **funct7**=0000101.

The **funct7**=0000101 block contains **clmul[hr]**, **min[u]**, and **max[u]**.

The encoding of **clmul**, **clmulr**, **clmulh** is identical to the encoding of **mulh**, **mulhsu**, **mulhu**, except that **op**[27]=1.

The encoding of **min[u]**/**max[u]** uses **funct3**=100..111. The **funct3** encoding matches **op**[31:29] of the AMO **min**/**max** functions.

The remaining instructions are encoded within **funct7**=0000100. The shift-like **shfl/unshfl** instructions uses the same **funct3** values as the shift operations. **bdep** and **bext** are encoded in a way so that **funct3**[2] selects the “direction”, similar to shift operations.

bmat[x] or use **funct3**=011 and **funct3**=111 in **funct7**=0000100.

pack occupies **funct3**=100 in **funct7**=0000100.

addwu and **subwu** are encoded like **addw** and **subw**, except that **op**[25]=1 and **op**[27]=1.

addu.w and **subu.w** are encoded like **addw** and **subw**, except that **op**[27]=1.

addiwu is encoded using **funct3**=100 (XOR) instead of **funct3**=000 in OP-32.

Finally, RV64 has **W* instructions for all bitmanip instructions, with the following exceptions:

andn, **cmix**, **cmov**, **min[u]**, **max[u]** have no **W* variants because they already behave in the way a **W* instruction would when presented with sign-extended 32-bit arguments.

bmatflip, **bmatxor**, **bmator** have no **W* variants because they are 64-bit only instructions.

crc32.[bhwd], **crc32c.[bhwd]** have no **W* variants because **crc32[c].w** is deemed sufficient.

There is no **[un]shfliw**, as a perfect outer shuffle always preserves the MSB bit, thus **[un]shfli** preserves proper sign extension when the upper bit in the control word is set. There’s still **[un]shflw** that masks that upper control bit and sign-extends the output.

Relevant instruction encodings from the base ISA are included in the table below and are marked with a ***.

3										2										1															
1	0	9	8	7	6	5	4	3	2	1	0	9	8	7	6	5	4	3	2	1	0	9	8	7	6	5	4	3	2	1	0				

funct7										rs2					rs1					f3					rd					opcode					R-type
rs3					f2					rs2					rs1					f3					rd					opcode					R4-type
imm															rs1					f3					rd					opcode					I-type
=====																																			
0000000										rs2					rs1					111					rd					0110011					AND*
0000000										rs2					rs1					110					rd					0110011					OR*
0000000										rs2					rs1					100					rd					0110011					XOR*
0100000										rs2					rs1					111					rd					0110011					ANDN
0100000										rs2					rs1					110					rd					0110011					ORN
0100000										rs2					rs1					100					rd					0110011					XNOR

0000000										rs2					rs1					001					rd					0110011					SLL*
0000000										rs2					rs1					101					rd					0110011					SRL*
0100000										rs2					rs1					101					rd					0110011					SRA*
0010000										rs2					rs1					001					rd					0110011					SLO
0010000										rs2					rs1					101					rd					0110011					SRO
0110000										rs2					rs1					001					rd					0110011					ROL
0110000										rs2					rs1					101					rd					0110011					ROR

0100100										rs2					rs1					001					rd					0110011					SBCLR
0010100										rs2					rs1					001					rd					0110011					SBSET
0110100										rs2					rs1					001					rd					0110011					SBINV
0100100										rs2					rs1					101					rd					0110011					SBEXT
0010100										rs2					rs1					101					rd					0110011					GORC
0110100										rs2					rs1					101					rd					0110011					GREV

00000					imm					rs1					001					rd					0010011					SLLI*					
00000					imm					rs1					101					rd					0010011					SRLI*					
01000					imm					rs1					101					rd					0010011					SRAI*					
00100					imm					rs1					001					rd					0010011					SLOI					
00100					imm					rs1					101					rd					0010011					SROI					
01100					imm					rs1					101					rd					0010011					RORI					

01001					imm					rs1					001					rd					0010011					SBCLRI					
00101					imm					rs1					001					rd					0010011					SBSETI					
01101					imm					rs1					001					rd					0010011					SBINVI					
01001					imm					rs1					101					rd					0010011					SBEXTI					
00101					imm					rs1					101					rd					0010011					GORCI					
01101					imm					rs1					101					rd					0010011					GREVI					

rs3					11					rs2					rs1					001					rd					0110011					CMIX
rs3					11					rs2					rs1					101					rd					0110011					CMOV
rs3					10					rs2					rs1					001					rd					0110011					FSL
rs3					10					rs2					rs1					101					rd					0110011					FSR
rs3					1					imm					rs1					101					rd					0010011					FSRI

3										2										1																	
1	0	9	8	7	6	5	4	3	2	1	0	9	8	7	6	5	4	3	2	1	0	9	8	7	6	5	4	3	2	1	0						

funct7										rs2					rs1					f3					rd					opcode							R-type
rs3					f2					rs2					rs1					f3					rd					opcode							R4-type
imm															rs1					f3					rd					opcode							I-type
=====																																					
0110000										00000					rs1					001					rd					0010011							CLZ
0110000										00001					rs1					001					rd					0010011							CTZ
0110000										00010					rs1					001					rd					0010011							PCNT
0110000										00011					rs1					001					rd					0010011							BMATFLIP

0110000										10000					rs1					001					rd					0010011							CRC32.B
0110000										10001					rs1					001					rd					0010011							CRC32.H
0110000										10010					rs1					001					rd					0010011							CRC32.W
0110000										10011					rs1					001					rd					0010011							CRC32.D
0110000										11000					rs1					001					rd					0010011							CRC32C.B
0110000										11001					rs1					001					rd					0010011							CRC32C.H
0110000										11010					rs1					001					rd					0010011							CRC32C.W
0110000										11011					rs1					001					rd					0010011							CRC32C.D

0000101										rs2					rs1					001					rd					0110011							CLMUL
0000101										rs2					rs1					010					rd					0110011							CLMULR
0000101										rs2					rs1					011					rd					0110011							CLMULH
0000101										rs2					rs1					100					rd					0110011							MIN
0000101										rs2					rs1					101					rd					0110011							MAX
0000101										rs2					rs1					110					rd					0110011							MINU
0000101										rs2					rs1					111					rd					0110011							MAXU

0000100										rs2					rs1					001					rd					0110011							SHFL
0000100										rs2					rs1					101					rd					0110011							UNSHFL
0000100										rs2					rs1					010					rd					0110011							BDEP
0000100										rs2					rs1					110					rd					0110011							BEXT
0000100										rs2					rs1					100					rd					0110011							PACK
0000100										rs2					rs1					011					rd					0110011							BMATOR
0100100										rs2					rs1					011					rd					0110011							BMATXOR
0000100										rs2					rs1					111					rd					0110011							BFP

0000010										imm					rs1					001					rd					0010011							SHFLI
0000010										imm					rs1					101					rd					0010011							UNSHFLI
=====																																					
immediate															rs1					000					rd					0011011							ADDIW*
immediate															rs1					100					rd					0011011							ADDIWU
00001					imm										rs1					001					rd					0011011							SLLIU.W

0000000										rs2					rs1					000					rd					0111011							ADDW*
0100000										rs2					rs1					000					rd					0111011							SUBW*
0000101										rs2					rs1					000					rd					0111011							ADDWU
0100101										rs2					rs1					000					rd					0111011							SUBWU
0000100										rs2					rs1					000					rd					0111011							ADDU.W
0100100										rs2					rs1					000					rd					0111011							SUBU.W

2.12 Future compressed instructions

The RISC-V ISA has no dedicated instructions for bitwise inverse (**not**). Instead **not** is implemented as **xori rd, rs, -1** and **neg** is implemented as **sub rd, x0, rs**.

In bitmanipulation code **not** is a very common operation. But there is no compressed encoding for those operation because there is no **c.xori** instruction.

On RV64 (and RV128) **zext.w** and **zext.d** (**pack** and **packw**) are commonly used to zero-extend unsigned values $< \text{XLEN}$.

It presumably would make sense for a future revision of the “C” extension to include compressed opcodes for those instructions.

An encoding with the constraint **rd = rs** would fit nicely in the reserved space in **c.addi16sp/c.lui**.

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0	
011			nzimm[9]				2					nzimm[4 6 8:7 5]		01		C.ADDI16SP (<i>RES, nzimm=0</i>)
011			nzimm[17]				rd≠{0,2}					nzimm[16:12]		01		C.LUI (<i>RES, nzimm=0; HINT, rd=0</i>)
011			0			00	rs1'/rd'					0		01		C.NOT
011			0			01	rs1'/rd'					0		01		C.ZEXT.W (<i>RV64/128</i>)
011			0			11	rs1'/rd'					0		01		C.ZEXT.D (<i>RV128</i>)

The entire RVC encoding space is 15.585 bits wide, the remaining reserved encoding space in RVC is 11.155 bits wide, not including space that is only reserved on RV32/RV64. This means that above encoding would use 0.0065% of the RVC encoding space, or 1.4% of the remaining reserved RVC encoding space. Preliminary experiments have shown that NOT instructions alone make up approximately 1% of bitmanipulation code size. [26]

2.13 Micro architectural considerations and macro-op fusion for bit-manipulation

2.13.1 Fast MUL, MULH, MULHSU, MULHU

A lot of bit manipulation code depends on “multiply with magic number”-tricks. Often those tricks need the upper half of the $2 \cdot \text{XLEN}$ product. Therefore decent performance for the **MUL** and especially **MULH[[S]U]** instructions is important for fast bit manipulation code.

2.13.2 Fused load-immediate sequences

Bit manipulation code, even more than other code, requires a lot of “magic numbers”, bitmasks, and other (usually large) constants. On some microarchitectures those can easily be loaded from a

nearby data section using load instructions. On other microarchitectures however this comes at a high cost, and it is more efficient to load immediates using a sequence of instructions.

Loading a 32-bit constant:

```
lui rd, imm
addi rd, rd, imm
```

On RV64 a 64 bit constant can be loaded by loading two 32-bit constants and combining them with a `PACK` instruction:

```
lui tmp, imm
addi tmp, tmp, imm
lui rd, imm
addi rd, rd, imm
pack rd, rd, tmp
```

(Without the temporary register and without the `PACK` instruction more complex/diverse sequences are used to load 64-bit immediates. But the `PACK` instruction streamlines the pattern and thus simplifies macro-op fusion.)

A 32-bit core should be capable of fusing the `lui+addi` pattern.

In addition to that, a 64 bit core may consider fusing the following sequences as well:

```
lui rd, imm
addi rd, rd, imm
pack rd, rd, rs2

lui rd, imm
pack rd, rd, rs2

addi rd, zero, imm
pack rd, rd, rs2
```

Furthermore, a core may consider fusing 32-bit immediate loads with any ALU instruction, not just `pack`:

```
lui rd, imm
addi rd, rd, imm
alu_op rd, rd, rs2

lui rd, imm
alu_op rd, rd, rs2

addi rd, zero, imm
alu_op rd, rd, rs2
```

And finally, a 64-bit core should fuse sequences with `addiwu` as well as `addi`, for loading unsigned 32-bit numbers that have their MSB set. This is often the case with masks in bit manipulation code.

2.13.3 Fused *-bfp sequences

The `bfp` instruction is most commonly used in one of the following sequences:

```
addi rd, zero, ...
pack rd, rs2, rd
bfp rd, rs1, rd

lui rd, ...
addi rd, rd, ...
bfp rd, rs1, rd
```

Either sequence only reads at most two registers and only writes one register, making them ideal candidates for macro-op fusion.

2.13.4 Fused *-not sequences

Preliminary experiments have shown that NOT instructions make up approximately 1% of bitmanipulation code size, more when looking at dynamic instruction count. [26]

Therefore it makes sense to fuse NOT instructions with other ALU instructions, if possible.

The most important form of NOT fusion is postfix fusion:

```
alu_op rd, rs1, rs2
not rd, rd
```

A future compressed NOT instruction would help keeping those fused sequences short.

2.13.5 Fused *-srli and *-srai sequences

Pairs of left and right shifts are common operations for extracting a bit field.

To extract the contiguous bit field starting at `pos` with length `len` from `rs` (with `pos > 0`, `len > 0`, and `pos + len ≤ XLEN`):

```
slli rd, rs, (XLEN-len-pos)
srli rd, rd, (XLEN-len)
```

Using `srai` instead of `srli` will sign-extend the extracted bit-field.

Similarly, placing a bit field with length `len` at the position `pos`:

```
slli rd, rs, (XLEN-len-pos)
srli rd, rd, (XLEN-len)
```

If possible, an implementation should fuse the following macro ops:

```
alu_op rd, rs1, rs2
srli rd, rd, imm
```

```
alu_op rd, rs1, rs2
srai rd, rd, imm
```

Note that the postfix right shift instruction can use a compressed encoding, yielding a 48-bit fused instruction if `alu_op` is a 32-bit instruction.

For generating masks, i.e. constants with one continuous run of 1 bits, a sequence like the following can be used that would utilize postfix fusion of right shifts:

```
sroi rd, zero, len
c.srli rd, (XLEN-len-pos)
```

This can be a useful sequence on RV64, where loading an arbitrary 64-bit constant would usually require at least 96 bits (using `c.ld`).

2.13.6 Fused sequences for logic operations

RISC-V has dedicated instructions for branching on equal/not-equal. But C code such as the following would require set-equal and set-not-equal instructions, similar to `slt`.

```
int is_equal = (a == b);
int is_noteq = (c != d);
```

Those can be implemented using the following fuse-able sequences:

```
sub rd, rs1, rs2
sltui rd, rd, 1
```

```
sub rd, rs1, rs2
sltu rd, zero, rd
```

Likewise for logic OR:

```
int logic_or = (c || d);
```

```
or rd, rs1, rs2
sltu rd, zero, rd
```

And for logic AND, if `rd == rs1`:

```
int logic_and = (c && d);
```

```
beq rd, zero, skip_sltu
sltu rd, zero, rs2
skip_sltu:
```

Note that the first instruction can be compressed in all four cases if `rd == rs1`.

2.13.7 Fused ternary ALU sequences

Architectures with support for ternary operations may want to support fusing two ALU operations.

```
alu_op rd, ...
alu_op rd, rd, ...
```

This would be a postfix-fusion pattern, extending the postfix shift-right fusion described in the previous section.

Candidates for this kind of postfix fusion would be simple ALU operations, specifically AND/OR/XOR/ADD/SUB and ANDI/ORI/XORI/ADDI/SUBI.

2.13.8 Pseudo-ops for fused sequences

Assembler pseudo-ops for not postfix fusion:

```
nand rd, rs1, rs2      ->  and rd, rs1, rs2; not rd, rd
nor  rd, rs1, rs2      ->  or  rd, rs1, rs2; not rd, rd
```

Assembler bitfield pseudo-ops for `sr[la]i` postfix fusion:

```
bfext rd, rs, len, pos  ->  slli rd, rs, (XLEN-len-pos); srai rd, rd, (XLEN-len)
bfextu rd, rs, len, pos ->  slli rd, rs, (XLEN-len-pos); srli rd, rd, (XLEN-len)
bfmak rd, len, pos      ->  sroi rd, zero, len; srli rd, rd, (XLEN-len-pos)
```

The names `bfext`, `bfextu`, and `bfmak` are borrowed from m88k, that had dedicated instructions of those names (without `bf`-prefix) with equivalent semantics. [3, p. 3-28]

Sign-extending bytes and half-words are special cases of `bfext`:

```
sext.b rd, rs  ->  slli rd, rs, (XLEN-8); srai rd, rd, (XLEN-8)
sext.h rd, rs  ->  slli rd, rs, (XLEN-16); srai rd, rd, (XLEN-16)
```

2.14 C intrinsics via `<rvintrin.h>`

A C header file `<rvintrin.h>` is provided that contains assembler templates for directly creating assembler instructions from C code.

The header defines `_rv_*(...)` functions that operate on the `long` data type, `_rv32_*(...)` functions that operate on the `int32_t` data type, and `_rv64_*(...)` functions that operate on the `int64_t` data type. The `_rv64_*(...)` functions are only available on RV64. See table 2.7 for a complete list of intrinsics defined in `<rvintrin.h>`.

Usage example:

```
#include <rvintrin.h>

int find_nth_set_bit(unsigned int value, int cnt) {
    return _rv32_ctz(_rv32_bdep(1 << cnt, value));
}
```

Defining `RVINTRIN_EMULATE` before including `<rvintrin.h>` will define plain C functions that emulate the behavior of the RISC-V instructions. This is useful for testing software on non-RISC-V platforms.

Instruction	RV32		RV64		
	<code>_rv_*</code>	<code>_rv32_*</code>	<code>_rv_*</code>	<code>_rv32_*</code>	<code>_rv64_*</code>
<code>clz</code>	✓	✓	✓	✓	✓
<code>ctz</code>	✓	✓	✓	✓	✓
<code>pcnt</code>	✓	✓	✓	✓	✓
<code>pack</code>	✓	✓	✓	✓	✓
<code>min</code>	✓	✓	✓	✓	✓
<code>minu</code>	✓	✓	✓	✓	✓
<code>max</code>	✓	✓	✓	✓	✓
<code>maxu</code>	✓	✓	✓	✓	✓
<code>sbset</code>	✓	✓	✓	✓	✓
<code>sbclr</code>	✓	✓	✓	✓	✓
<code>sbinv</code>	✓	✓	✓	✓	✓
<code>sbext</code>	✓	✓	✓	✓	✓
<code>sll</code>	✓	✓	✓	✓	✓
<code>srl</code>	✓	✓	✓	✓	✓
<code>sra</code>	✓	✓	✓	✓	✓
<code>slo</code>	✓	✓	✓	✓	✓
<code>sro</code>	✓	✓	✓	✓	✓
<code>rol</code>	✓	✓	✓	✓	✓
<code>ror</code>	✓	✓	✓	✓	✓
<code>grev</code>	✓	✓	✓	✓	✓
<code>gorc</code>	✓	✓	✓	✓	✓
<code>shfl</code>	✓	✓	✓	✓	✓
<code>unshfl</code>	✓	✓	✓	✓	✓
<code>bfp</code>	✓	✓	✓	✓	✓
<code>bext</code>	✓	✓	✓	✓	✓
<code>bdep</code>	✓	✓	✓	✓	✓
<code>clmul</code>	✓	✓	✓	✓	✓
<code>clmulh</code>	✓	✓	✓	✓	✓
<code>clmulr</code>	✓	✓	✓	✓	✓
<code>bmatflip</code>			✓		✓
<code>bmator</code>			✓		✓
<code>bmatxor</code>			✓		✓
<code>fsl</code>	✓	✓	✓	✓	✓
<code>fsr</code>	✓	✓	✓	✓	✓
<code>cmix</code>	✓		✓		
<code>cmov</code>	✓		✓		
<code>crc32_b</code>	✓		✓		
<code>crc32_h</code>	✓		✓		
<code>crc32_w</code>	✓		✓		
<code>crc32_d</code>			✓		
<code>crc32c_b</code>	✓		✓		
<code>crc32c_h</code>	✓		✓		
<code>crc32c_w</code>	✓		✓		
<code>crc32c_d</code>			✓		

Table 2.7: C intrinsics defined in `<rvintrin.h>`

Chapter 3

Reference Implementations

3.1 Verilog reference implementations

We have implemented Verilog cores for all instructions proposed in this specification. These cores are permissively licensed under the ISC license and can be obtained from <https://github.com/riscv/riscv-bitmanip/tree/master/verilog>.

For evaluation purposes we synthesized these cores for RV32 and RV64 to the following mockup ASIC cell library:

Cell	Gate Count	Cell	Gate Count
NOT	0.5	AOI3	1.5
NAND	1	OAI3	1.5
NOR	1	AOI4	2
XOR	3	OAI4	2
XNOR	3	NMUX	2.5
DFF	4	MUX	3

For comparison we also synthesized the rocket-chip MulDiv cores obtained using the following rocket-chip configurations:

```
class MulDivConfig64 extends Config(  
    new WithFastMulDiv ++  
    new DefaultConfig  
)  
  
class MulDivConfig32 extends Config(  
    new WithRV32 ++  
    new WithFastMulDiv ++  
    new DefaultConfig  
)
```




Figure 3.1: Area of 32-bit Rocket MulDiv core (left) compared to a complete implementation of all 32-bit instructions proposed in this specification except CRC instructions (right).

The following table lists the verilog reference cores and the instructions they implement:

Module	Instructions
<code>rvb_bextdep</code>	<code>bext bdep grev gorc shfl unshfl</code>
<code>rvb_clmul</code>	<code>clmul clmulr clmulh</code>
<code>rvb_shifter</code>	<code>sll srl sra slo sro rol ror fsl fsr slliu.w</code> <code>sbset sbclr sbinv sbext bfp</code>
<code>rvb_bmatxor</code>	<code>bmatxor bmator</code>
<code>rvb_simple</code>	<code>min max minu maxu andn orn</code> <code>xnor pack cmix cmov addiwu addwu</code> <code>subwu adduw subuw</code>
<code>rvb_bitcnt</code>	<code>clz ctz pcnt bmatflip</code>
<code>rvb_full</code>	All of the above

On RV64 these cores also implement all `*W` instruction variants of the above instructions.

Note that `rvb_shifter` also implements the base ISA `sll`, `srl`, and `sra` instructions. Thus it can replace an existing implementation of the base ISA shift instructions.

Fig. 3.1 shows the area comparison for RV32 and fig. 3.2 shows the comparison for RV64. The area of the red frame surrounding the blue `rvb.*` modules accurately represents the added area by the `rvb_full` wrapper module.

Regarding timing we evaluate the longest paths for `rvb_full` and rocket-chip MulDiv, measured in gate delays:

	RV32	RV64
<code>rvb_full</code>	30	57
MulDiv	43	68




Figure 3.2: Area of 64-bit Rocket MulDiv core (left) compared to a complete implementation of all 64-bit instructions proposed in this specification except CRC instructions (right).

All `rvb.*` reference cores provide single-cycle implementations of their functions, with the exception of `rvb.clmul` which requires 4 cycles for a 32-bit carry-less multiply and 8 cycles for a 64-bit carry-less multiply.

3.2 Fast C reference implementations

GCC has intrinsics for the bit counting instructions `clz`, `ctz`, and `pcnt`. So a performance-sensitive application (such as an emulator) should probably just use those:

```
uint32_t fast_clz32(uint32_t rs1)
{
    if (rs1 == 0)
        return 32;
    assert(sizeof(int) == 4);
    return __builtin_clz(rs1);
}

uint64_t fast_clz64(uint64_t rs1)
{
    if (rs1 == 0)
        return 64;
    assert(sizeof(long long) == 8);
    return __builtin_clzll(rs1);
}
```

```

uint32_t fast_ctz32(uint32_t rs1)
{
    if (rs1 == 0)
        return 32;
    assert(sizeof(int) == 4);
    return __builtin_ctz(rs1);
}

uint64_t fast_ctz64(uint64_t rs1)
{
    if (rs1 == 0)
        return 64;
    assert(sizeof(long long) == 8);
    return __builtin_ctzll(rs1);
}

uint32_t fast_pcmt32(uint32_t rs1)
{
    assert(sizeof(int) == 4);
    return __builtin_popcount(rs1);
}

uint64_t fast_pcmt64(uint64_t rs1)
{
    assert(sizeof(long long) == 8);
    return __builtin_popcountll(rs1);
}

```

For processors with BMI2 support GCC has intrinsics for bit extract and bit deposit instructions (compile with `-mbmi2` and include `<x86intrin.h>`):

```

uint32_t fast_bext32(uint32_t rs1, uint32_t rs2)
{
    return _pext_u32(rs1, rs2);
}

uint64_t fast_bext64(uint64_t rs1, uint64_t rs2)
{
    return _pext_u64(rs1, rs2);
}

uint32_t fast_bdep32(uint32_t rs1, uint32_t rs2)
{
    return _pdep_u32(rs1, rs2);
}

uint64_t fast_bdep64(uint64_t rs1, uint64_t rs2)
{
    return _pdep_u64(rs1, rs2);
}

```

For other processors we need to provide our own implementations. The following implementation

is a good compromise between code complexity and runtime:

```
uint_xlen_t fast_bext(uint_xlen_t rs1, uint_xlen_t rs2)
{
    uint_xlen_t c = 0, i = 0, mask = rs2;
    while (mask) {
        uint_xlen_t b = mask & ~((mask | (mask-1)) + 1);
        c |= (rs1 & b) >> (fast_ctz(b) - i);
        i += fast_pcmt(b);
        mask -= b;
    }
    return c;
}

uint_xlen_t fast_bdep(uint_xlen_t rs1, uint_xlen_t rs2)
{
    uint_xlen_t c = 0, i = 0, mask = rs2;
    while (mask) {
        uint_xlen_t b = mask & ~((mask | (mask-1)) + 1);
        c |= (rs1 << (fast_ctz(b) - i)) & b;
        i += fast_pcmt(b);
        mask -= b;
    }
    return c;
}
```

For the other Bitmanip instructions the C reference functions given in Chapter 2 are already reasonably efficient.

Chapter 4

Evaluation

This chapter contains a collection of short code snippets and algorithms using the Bitmanip extension for evaluation purposes. For the sake of simplicity we assume RV32 for most examples in this chapter.

4.1 Basic Bitmanipulation

4.1.1 Bitfield extract

Extracting a bit field of length `len` at position `pos` can be done using two shift operations.

```
slli a0, a0, (XLEN-len-pos)
srli a0, a0, (XLEN-len)
```

Or using `srai` for a signed bit-field.

```
slli a0, a0, (XLEN-len-pos)
srai a0, a0, (XLEN-len)
```

4.1.2 Parity check

The parity of a word (xor of all bits) is the LSB of the population count.

```
pcnt a0, a0
andi a0, a0, 1
```

4.1.3 Rank and select

Rank and select are fundamental operations in succinct data structures [19].

`select(a0, a1)` returns the position of the `a1`th set bit in `a0`. It can be implemented efficiently using `bdep` and `ctz`:

```
select:
    li a2, 1
    sll a1, a2, a1
    bdep a0, a1, a0
    ctz a0, a0
    ret
```

`rank(a0, a1)` returns the number of set bits in `a0` up to and including position `a1`.

```
rank:
    not a1, a1
    sll a0, a1
    pcnt a0, a0
    ret
```

4.1.4 Packing bytes in word

The following code packs the lower 8 bits from `a0`, `a1`, `a2`, `a3` into a 32-bit word returned in `a0`, ignoring other bits in the input values.

```
pack a0, a0, a1
pack a1, a2, a3
shfl a0, a0, 8
shfl a1, a1, 8
pack a0, a0, a1
```

This replaces either 4 store-byte instructions followed by one load-word instruction, or something like the following sequence.

```
andi a0, a0, 255
andi a1, a1, 255
andi a2, a2, 255
pack a0, a0, a2
pack a1, a1, a3
slli a1, a1, 8
or a0, a0, a1
```

4.1.5 Counting trailing non-zero bytes

Counting the trailing (LSB-end) non-zero bytes in a word is a helpful operation in optimized implementations of `strlen()` and `strcpy()`:

```
int count_trailing_nonzero_bytes(long x)
{
    return _rv_ctz(~_rv_orc_b(x)) >> 3;
}
```

4.1.6 Finding bytes of certain values

Finding zero bytes is a useful operations for `strchr()` and `memchr()`:

```
bool check_zero_bytes(long x)
{
    return ~_rv_orc_b(x) != 0;
}
```

To find other bytes we simply XOR the value with a mask of the byte value we are looking for:

```
bool check_byte(long x, unsigned char c)
{
    return ~_rv_orc_b(x ^ _rv_orc8(c)) != 0;
}
```

These schemes can easily be extended with `ctz` and `pcnt` to perform operations such as counting the number of bytes of a certain value within a word, or finding the position of the first such byte.

4.1.7 Fill right of most significant set bit

The “fill right” or “fold right” operation is a pattern commonly used in bit manipulation code. [8]

The straight-forward RV64 implementation requires 12 instructions:

```
uint64_t rfill(uint64_t x)
{
    x |= x >> 1;    // SRLI, OR
    x |= x >> 2;    // SRLI, OR
    x |= x >> 4;    // SRLI, OR
    x |= x >> 8;    // SRLI, OR
    x |= x >> 16;   // SRLI, OR
    x |= x >> 32;   // SRLI, OR
    return x;
}
```

With `clz` it can be implemented in only 4 instructions. Notice the handling of the case where `x=0` using `sltiu+addi`.

```

uint64_t rfill_clz(uint64_t x)
{
    uint64_t t;
    t = clz(x);          // CLZ
    x = (!x)-1;          // SLTIU, ADDI
    x = x >> (t & 63);    // SRL
    return x;
}

```

Alternatively, a Trailing Bit Manipulation (TBM) code pattern can be used together with `rev` to implement this function in 4 instructions:

```

uint64_t rfill_rev(uint64_t x)
{
    x = rev(x);          // GREVI
    x = x | ~(x - 1);    // ADDI, ORN
    x = rev(x);          // GREVI
    return x;
}

```

Finally, there is another implementation in 4 instructions using BMATOR, if we do not count the extra instructions for loading utility matrices.

```

uint64_t rfill_bmat(uint64_t x)
{
    uint64_t m0, m1, m2, t;

    m0 = 0xFF7F3F1F0F070301LL; // LD
    m1 = bmatflip(m0 << 8);    // SLLI, BMATFLIP
    m2 = -1LL;                  // ADDI

    t = bmator(x, m0);         // BMATOR
    x = bmator(x, m2);         // BMATOR
    x = bmator(m1, x);         // BMATOR
    x |= t;                    // OR

    return x;
}

```

4.1.8 Round to next power of two

One common application of `rfill()` is rounding up to the next power of two:

```

uint64_t round_pow2(uint64_t x)
{
    return rfill(x-1)+1;
}

```

This can also be implemented in just 4 instructions, if we don't care about the case where the above code overflows because `x` is already larger than the largest power-of-two representable in an


```
uint64_t
uint64_t round_pow2(uint64_t x)
{
    uint64_t t;
    t = clz(x-1);    // ADDI, CLZ
    x = ror(!!x, t); // SLTU, ROR
    return x;
}
```

Note that this code handles $0 \rightarrow 0$ and $1 \rightarrow 1$ correctly, i.e. equivalent to `rfill(x-1)+1`.

4.2 Funnel shifts

A funnel shift takes two XLEN registers, concatenates them to a $2 \times \text{XLEN}$ word, shifts that by a certain amount, then returns the lower half of the result for a right shift and the upper half of the result for a left shift.

The `fsl`, `fsr`, and `fsri` instructions perform funnel shifts.

4.2.1 Bigint shift

A common application for funnel shifts is shift operations in bigint libraries.

For example, the following functions implement rotate-shift operations for bigints made from `n` XLEN words.

```

void bigint_rol(uint_xlen_t data[], int n, int shamt)
{
    if (n <= 0)
        return;

    uint_xlen_t buffer = data[n-1];
    for (int i = n-1; i > 0; i--)
        data[i] = fsl(data[i], shamt, data[i-1]);
    data[0] = fsl(data[0], shamt, buffer);
}

void bigint_ror(uint_xlen_t data[], int n, int shamt)
{
    if (n <= 0)
        return;

    uint_xlen_t buffer = data[0];
    for (int i = 0; i < n-1; i++)
        data[i] = fsr(data[i], shamt, data[i+1]);
    data[n-1] = fsr(data[n-1], shamt, buffer);
}

```

These version only works for shift-amounts $< \text{XLEN}$. But functions supporting other kinds of shift operations, or shifts $\geq \text{XLEN}$ can easily be built with `fsl` and `fsr`.

4.2.2 Parsing bit-streams

The following function parses `n` 27-bit words from a packed array of `XLEN` words:

```

void parse_27bit(uint_xlen_t *idata, uint_xlen_t *odata, int n)
{
    uint_xlen_t lower = 0, upper = 0;
    int reserve = 0;

    while (n--) {
        if (reserve < 27) {
            uint_xlen_t buf = *(idata++);
            lower |= sll(buf, reserve);
            upper = reserve ? srl(buf, -reserve) : 0;
            reserve += XLEN;
        }
        *(odata++) = lower & ((1 << 27)-1);
        lower = fsr(lower, 27, upper);
        upper = srl(upper, 27);
        reserve -= 27;
    }
}

```

And here the same thing in RISC-V assembler:

```

parse_27bit:
    li t1, 0          ; lower
    li t2, 0          ; upper
    li t3, 0          ; reserve
    li t4, 27         ; shamt
    slo t5, zero, t4   ; mask
    beqz a2, endloop   ; while (n--)
loop:
    addi a2, a2, -1
    bge t3, t4, output ; if (reserve < 27)
    lw t6, 0(a0)        ; buf = *(idata++)
    addi a0, a0, 4
    sll t7, t6, t3       ; lower |= sll(buf, reserve)
    or t1, t1, t7
    sub t7, zero, t3     ; upper = reserve ? srl(buf, -reserve) : 0
    srl t7, t6, t7
    cmov t2, t3, t7, zero
    addi t3, t3, 32      ; reserve += XLEN;
output:
    and t6, t1, t5       ; *(odata++) = lower & ((1 << 27)-1)
    sw t6, 0(a1)
    addi a1, a1, 4
    fsr t1, t1, t2, t4   ; lower = fsr(lower, 27, upper)
    srl t2, t2, t4       ; upper = srl(upper, 27)
    sub t3, t3, t4       ; reserve -= 27
    bnez a2, loop        ; while (n--)
endloop:
    ret

```

A loop iteration without fetch is 9 instructions long, and a loop iteration with fetch is 17 instructions long.

Without ternary operators that would be 13 instructions and 22 instructions, i.e. assuming one cycle per instruction, that function would be about 30% slower without ternary instructions.

4.2.3 Fixed-point multiply

A fixed-point multiply is simply an integer multiply, followed by a right shift. If the entire dynamic range of XLEN bits should be useable for the factors, then the product before shift must be $2 \times \text{XLEN}$ wide. Therefore `mul+mulh` is needed for the multiplication, and funnel shift instructions can help with the final right shift. For fixed-point numbers with N fraction bits:

```

mul_fracN:
    mulh a2, a0, a1
    mul a0, a0, a1
    fsri a0, a0, a2, N
    ret

```

4.3 Arbitrary bit permutations

This section lists code snippets for computing arbitrary bit permutations that are defined by data (as opposed to bit permutations that are known at compile time and can likely be compiled into shift-and-mask operations and/or a few instances of bext/bdep).

4.3.1 Using butterfly operations

The following macro performs a stage-N butterfly operation on the word in `a0` using the mask in `a1`.

```
grevi a2, a0, (1 << N)
cmix a0, a1, a2, a0
```

The bitmask in `a1` must be preformatted correctly for the selected butterfly stage. A butterfly operation only has a $XLEN/2$ wide control word. The following macros format the mask assuming those $XLEN/2$ bits in the lower half of `a1` on entry:

```
bfly_msk_0:
    pack a1, a1, a1
    zip a1, a1
```

```
bfly_msk_1:
    pack a1, a1, a1
    zip2 a1, a1
```

```
bfly_msk_2:
    pack a1, a1, a1
    zip4 a1, a1
```

```
...
```

A sequence of $2 \cdot \log_2(XLEN) - 1$ butterfly operations can perform any arbitrary bit permutation (Beneš network):

```
butterfly(LOG2_XLEN-1)
butterfly(LOG2_XLEN-2)
...
butterfly(0)
...
butterfly(LOG2_XLEN-2)
butterfly(LOG2_XLEN-1)
```

Many permutations arising from real-world applications can be implemented using shorter sequences. For example, any sheep-and-goats operation (SAG, see section 4.3.4) with either the sheep or the goats bit reversed can be implemented in $\log_2(\text{XLEN})$ butterfly operations.

Reversing a permutation implemented using butterfly operations is as simple as reversing the order of butterfly operations.

4.3.2 Using omega-flip networks

The omega operation is a stage-0 butterfly preceded by a zip operation:

```
zip a0, a0
grevi a2, a0, 1
cmix a0, a1, a2, a0
```

The flip operation is a stage-0 butterfly followed by an unzip operation:

```
grevi a2, a0, 1
cmix a0, a1, a2, a0
unzip a0, a0
```

A sequence of $\log_2(\text{XLEN})$ omega operations followed by $\log_2(\text{XLEN})$ flip operations can implement any arbitrary 32 bit permutation.

As for butterfly networks, permutations arising from real-world applications can often be implemented using a shorter sequence.

4.3.3 Using baseline networks

Another way of implementing arbitrary 32 bit permutations is using a baseline network followed by an inverse baseline network.

A baseline network is a sequence of $\log_2(\text{XLEN})$ butterfly(0) operations interleaved with unzip operations. For example, a 32-bit baseline network:

```
butterfly(0)
unzip
butterfly(0)
unzip.h
butterfly(0)
unzip.b
butterfly(0)
unzip.n
butterfly(0)
```

An inverse baseline network is a sequence of $\log_2(\text{XLEN})$ butterfly(0) operations interleaved with zip operations. The order is opposite to the order in a baseline network. For example, a 32-bit inverse baseline network:

```
butterfly(0)
zip.n
butterfly(0)
zip.b
butterfly(0)
zip.h
butterfly(0)
zip
butterfly(0)
```

A baseline network followed by an inverse baseline network can implement any arbitrary bit permutation.

4.3.4 Using sheep-and-goats

The Sheep-and-goats (SAG) operation is a common operation for bit permutations. It moves all the bits selected by a mask (goats) to the LSB end of the word and all the remaining bits (sheep) to the MSB end of the word, without changing the order of sheep or goats.

The SAG operation can easily be performed using `bext` (data in `a0` and mask in `a1`):

```
bext a2, a0, a1
not a1, a1
bext a0, a0, a1
pcnt a1, a1
ror a0, a0, a1
or a0, a0, a2
```

Any arbitrary bit permutation can be implemented in $\log_2(\text{XLEN})$ SAG operations.

The Hacker's Delight describes an optimized standard C implementation of the SAG operation. Their algorithm takes 254 instructions (for 32 bit) or 340 instructions (for 64 bit) on their reference RISC instruction set. [9, p. 152f, 162f]

4.3.5 Using bit-matrix multiply

`bat[x]or` performs a permutation of bits within each byte when used with a permutation matrix in `rs2`, and performs a permutation of bytes when used with a permutation matrix in `rs1`.

4.4 Mirroring and rotating bitboards

Bitboards are 64-bit bitmasks that are used to represent part of the game state in chess engines (and other board game AIs). The bits in the bitmask correspond to squares on a 8×8 chess board:

```

56 57 58 59 60 61 62 63
48 49 50 51 52 53 54 55
40 41 42 43 44 45 46 47
32 33 34 35 36 37 38 39
24 25 26 27 28 29 30 31
16 17 18 19 20 21 22 23
 8  9 10 11 12 13 14 15
 0  1  2  3  4  5  6  7

```

Many bitboard operations are simple straight-forward operations such as bitwise-AND, but mirroring and rotating bitboards can take up to 20 instructions on x86.

4.4.1 Mirroring bitboards

Flipping horizontally or vertically can easily be done with `grevi`:

Flip horizontal:

63 62 61 60 59 58 57 56	RISC-V Bitmanip:
55 54 53 52 51 50 49 48	<code>rev.b</code>
47 46 45 44 43 42 41 40	
39 38 37 36 35 34 33 32	
31 30 29 28 27 26 25 24	x86:
23 22 21 20 19 18 17 16	13 operations
15 14 13 12 11 10 9 8	
7 6 5 4 3 2 1 0	

Flip vertical:

0 1 2 3 4 5 6 7	RISC-V Bitmanip:
8 9 10 11 12 13 14 15	<code>rev8</code>
16 17 18 19 20 21 22 23	
24 25 26 27 28 29 30 31	
32 33 34 35 36 37 38 39	x86:
40 41 42 43 44 45 46 47	<code>bswap</code>
48 49 50 51 52 53 54 55	
56 57 58 59 60 61 62 63	

Rotating by 180 (flip horizontal and vertical):

Rotate 180:

```

 7  6  5  4  3  2  1  0    RISC-V Bitmanip:
15 14 13 12 11 10  9  8      rev
23 22 21 20 19 18 17 16
31 30 29 28 27 26 25 24
39 38 37 36 35 34 33 32    x86:
47 46 45 44 43 42 41 40      14 operations
55 54 53 52 51 50 49 48
63 62 61 60 59 58 57 56

```

4.4.2 Rotating bitboards

Using `zip` a bitboard can be transposed easily:

```

Transpose:
 7 15 23 31 39 47 55 63    RISC-V Bitmanip:
 6 14 22 30 38 46 54 62      zip, zip, zip
 5 13 21 29 37 45 53 61
 4 12 20 28 36 44 52 60
 3 11 19 27 35 43 51 59    x86:
 2 10 18 26 34 42 50 58      18 operations
 1  9 17 25 33 41 49 57
 0  8 16 24 32 40 48 56

```

A rotation is simply the composition of a flip operation and a transpose operation. This takes 19 operations on x86 [2]. With Bitmanip the rotate operation only takes 4 operations:

```

rotate_bitboard:
    rev8 a0, a0
    zip a0, a0
    zip a0, a0
    zip a0, a0

```

4.4.3 Explanation

The bit indices for a 64-bit word are 6 bits wide. Let $i[5:0]$ be the index of a bit in the input, and let $i'[5:0]$ be the index of the same bit after the permutation.

As an example, a rotate left shift by N can be expressed using this notation as $i'[5:0] = i[5:0] + N \pmod{64}$.

The GREV operation with shamt N is $i'[5:0] = i[5:0] \text{ XOR } N$.

And a GZIP operation corresponds to a rotate left shift by one position of any contiguous region of $i[5:0]$. For example, `zip` is a left rotate shift of the entire bit index:

$$i'[5:0] = \{i[4:0], i[5]\}$$

And `zip4` performs a left rotate shift on bits 5:2:

$$i'[5:0] = \{i[4:2], i[5], i[1:0]\}$$

In a bitboard, `i[2:0]` corresponds to the X coordinate of a board position, and `i[5:3]` corresponds to the Y coordinate.

Therefore flipping the board horizontally is the same as negating bits `i[2:0]`, which is the operation performed by `grevi rd, rs, 7 (rev.b)`.

Likewise flipping the board vertically is done by `grevi rd, rs, 56 (rev8)`.

Finally, transposing corresponds by swapping the lower and upper half of `i[5:0]`, or rotate shifting `i[5:0]` by 3 positions. This can easily done by rotate shifting the entire `i[5:0]` by one bit position (`zip`) three times.

4.4.4 Rotating Bitcubes

Let's define a bitcube as a $4 \times 4 \times 4$ cube with $x = i[1:0]$, $y = i[3:2]$, and $z = i[5:4]$. Using the same methods as described above we can easily rotate a bitcube by 90° around the X-, Y-, and Z-axis:

<code>rotate_x:</code>	<code>rotate_y:</code>	<code>rotate_z:</code>
<code>rev16 a0, a0</code>	<code>rev.n a0, a0</code>	<code>rev4.h</code>
<code>zip4 a0, a0</code>	<code>zip a0, a0</code>	<code>zip.h a0, a0</code>
<code>zip4 a0, a0</code>	<code>zip a0, a0</code>	<code>zip.h a0, a0</code>
	<code>zip4 a0, a0</code>	
	<code>zip4 a0, a0</code>	

4.5 Inverting Xorshift RNGs

Xorshift RNGs are a class of fast RNGs for different bit widths. There are 648 Xorshift RNGs for 32 bits, but this is the one that the author of the original Xorshift RNG paper recommends. [18, p. 4]

```
uint32_t xorshift32(uint32_t x)
{
    x ^= x << 13;
    x ^= x >> 17;
    x ^= x << 5;
    return x;
}
```

This function of course has been designed and selected so it's efficient, even without special bit-manipulation instructions. So let's look at the inverse instead. First, the naïve form of inverting this function:

```
uint32_t xorshift32_inv(uint32_t x)
{
    uint32_t t;
    t = x ^ (x << 5);
    t = x ^ (t << 5);
    t = x ^ (t << 5);
    t = x ^ (t << 5);
    t = x ^ (t << 5);
    x = x ^ (t << 5);
    x = x ^ (x >> 17);
    t = x ^ (x << 13);
    x = x ^ (t << 13);
    return x;
}
```

This translates to 18 RISC-V instructions, not including the function call overhead.

Obviously the C expression $x \wedge (x \gg 17)$ is already its own inverse (because $17 \geq XLEN/2$) and therefore already has an efficient inverse. But the two other blocks can easily be implemented using a single `clmul` instruction each:

```
uint32_t xorshift32_inv(uint32_t x)
{
    x = clmul(x, 0x42108421);
    x = x ^ (x >> 17);
    x = clmul(x, 0x04002001);
    return x;
}
```

This are 8 RISC-V instructions, including 4 instructions for loading the constants, but not including the function call overhead.

An optimizing compiler could easily generate the `clmul` instructions and the magic constants from the C code for the naïve implementation. ($0x04002001 = (1 \ll 2*13) \mid (1 \ll 13) \mid 1$ and $0x42108421 = (1 \ll 6*5) \mid (1 \ll 5*5) \mid \dots \mid (1 \ll 5) \mid 1$)

The obvious remaining question is “if `clmul(x, 0x42108421)` is the inverse of $x \wedge (x \ll 5)$, what's the inverse of $x \wedge (x \gg 5)$?” It's `clmulr(x, 0x84210842)`, where $0x84210842$ is the bit-reversal of $0x42108421$.

A special case of `xorshift` is $x \wedge (x \gg 1)$, which is a gray encoder. The corresponding gray decoder is `clmulr(x, 0xffffffff)`.

4.6 Cyclic redundancy checks (CRC)

There are special instructions for performing CRCs using the two most widespread 32-bit CRC polynomials, CRC-32 and CRC-32C.

CRCs with other polynomials can be computed efficiently using CLMUL. The following examples are using CRC32Q.

The easiest way of implementing CRC32Q with clmul is using a Barrett reduction. On RV32:

```
uint32_t crc32q_simple(const uint32_t *data, int length)
{
    uint32_t P = 0x814141AB; // CRC polynomial (implicit x^32)
    uint32_t mu = 0xFEFF7F62; // x^64 divided by CRC polynomial
    uint32_t mu1 = 0xFF7FBFB1; // "mu" with leading 1, shifted right by 1 bit
    uint32_t crc = 0;

    for (int i = 0; i < length; i++) {
        crc ^= rev8(data[i]);
        crc = clmulr(crc, mu1);
        crc = clmul(crc, P);
    }

    return crc;
}
```

The following python code calculates the value of mu for a given CRC polynomial:

```
def polydiv(dividend, divisor):
    quotient = 0
    while dividend.bit_length() >= divisor.bit_length():
        i = dividend.bit_length() - divisor.bit_length()
        dividend = dividend ^ (divisor << i)
        quotient |= 1 << i
    return quotient

P = 0x814141AB
print("0x%X" % (polydiv(1<<64, P))) # prints 0xFEFF7F62
```

A more efficient method would be the following, which processes 64-bit at a time (RV64):

```

uint32_t crc32q_fast(const uint64_t *p, int len)
{
    uint64_t P = 0x1814141ABLL;    // CRC polynomial
    uint64_t k1 = 0xA1FA6BECLL;    // rest of  $x^{128}$  divided by CRC polynomial
    uint64_t k2 = 0x9BE9878FLL;    // rest of  $x^{96}$  divided by CRC polynomial
    uint64_t k3 = 0xB1EFC5F6LL;    // rest of  $x^{64}$  divided by CRC polynomial
    uint64_t mu = 0x1FEFF7F62LL;    //  $x^{64}$  divided by CRC polynomial

    uint64_t a0, a1, a2, t1, t2;

    assert(len >= 2);
    a0 = rev8(p[0]);
    a1 = rev8(p[1]);

    // Main loop: Reduce to 2x 64 bits

    for (const uint64_t *t0 = p+2; t0 != p+len; t0++)
    {
        a2 = rev8(*t0);
        t1 = clmulh(a0, k1);
        t2 = clmul(a0, k1);
        a0 = a1 ^ t1;
        a1 = a2 ^ t2;
    }

    // Reduce to 64 bit, add 32 bit zero padding

    t1 = clmulh(a0, k2);
    t2 = clmul(a0, k2);

    a0 = (a1 >> 32) ^ t1;
    a1 = (a1 << 32) ^ t2;

    t2 = clmul(a0, k3);
    a1 = a1 ^ t2;

    // Barrett Reduction

    t1 = clmul(a1 >> 32, mu);
    t2 = clmul(t1 >> 32, P);
    a0 = a1 ^ t2;

    return a0;
}

```

The main idea is to transform an array of arbitrary length to an array with the same CRC that's only two 64-bit elements long. (That's the “Main loop” portion of above code.)

Then we further reduce it to just 64-bit. And then we use a Barrett reduction to get the final 32-bit

result.

The following python code can be used to calculate the “magic constants” `k1`, `k2`, and `k3`:

```
def polymod(dividend, divisor):
    quotient = 0
    while dividend.bit_length() >= divisor.bit_length():
        i = dividend.bit_length() - divisor.bit_length()
        dividend = dividend ^ (divisor << i)
        quotient |= 1 << i
    return dividend

print("0x%X" % (polymod(1<<128, P))) # prints 0xA1FA6BEC
print("0x%X" % (polymod(1<< 96, P))) # prints 0x9BE9878F
print("0x%X" % (polymod(1<< 64, P))) # prints 0xB1EFC5F6
```

The above example code is taken from [25]. A more detailed descriptions of the algorithms employed can be found in [11].

4.7 Decoding RISC-V immediates

The following code snippets decode and sign-extend the immediate from RISC-V S-type, B-type, J-type, and CJ-type instructions. They are nice “nothing up my sleeve”-examples for real-world bit permutations.

31	27	26	25	24	20	19	15	14	12	11	7	6	0		
imm[11:5]										imm[4:0]				S-type	
imm[12 10:5]										imm[4:1 11]				B-type	
imm[20 10:1 11 19:12]															J-type

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0	
imm[11 4 9:8 10 6 7 3:1 5]																CJ-type

decode_s:

```
li t0, 0xfe000f80
bext a0, a0, t0
c.slli a0, 20
c.srai a0, 20
ret
```

decode_b:

```
li t0, 0xea800aa
rori a0, a0, 8
grevi a0, a0, 8
shfli a0, a0, 7
bext a0, a0, t0
c.slli a0, 20
c.srai a0, 19
ret
```

```

decode_j:
    li t0, 0x800003ff
    li t1, 0x800ff000
    bext a1, a0, t1
    c.slli a1, 23
    rori a0, a0, 21
    bext a0, a0, t0
    c.slli a0, 12
    c.or a0, a1
    c.srai a0, 11
    ret

// variant 1 (with RISC-V Bitmanip)
decode_cj:
    li t0, 0x28800001
    li t1, 0x000016b8
    li t2, 0xb4e00000
    li t3, 0x4b000000
    bext a1, a0, t1
    bdep a1, a1, t2
    rori a0, a0, 11
    bext a0, a0, t0
    bdep a0, a0, t3
    c.or a0, a1
    c.srai a0, 20
    ret

```

```

// variant 2 (without RISC-V Bitmanip)
decode_cj:
    srli a5, a0, 2
    srli a4, a0, 7
    c.andi a4, 16
    slli a3, a0, 3
    c.andi a5, 14
    c.add a5, a4
    andi a3, a3, 32
    srli a4, a0, 1
    c.add a5, a3
    andi a4, a4, 64
    slli a2, a0, 1
    c.add a5, a4
    andi a2, a2, 128
    srli a3, a0, 1
    slli a4, a0, 19
    c.add a5, a2
    andi a3, a3, 768
    c.slli a0, 2
    c.add a5, a3
    andi a0, a0, 1024
    c.srai a4, 31
    c.add a5, a0
    slli a0, a4, 11
    c.add a0, a5
    ret

```

Change History

Date	Rev	Changes
2017-07-17	0.10	Initial Draft
2017-11-02	0.11	Remove roli, assembler can convert it to use a rori Remove bitwise subset and replace with andc Doc source text same base for study and spec. Fix typos
2017-11-30	0.32	Jump rev number to be on par with associated Study Move pdep/pext into spec draft and called it scatter-gather
2018-04-07	0.33	Move to github, throw out study, convert from .md to .tex Fix typos and fix some reference C implementations Rename bgat/bsca to bext/bdep Remove post-add immediate from clz Clean up encoding tables and code sections
2018-04-20	0.34	Add GREV, CTZ, and compressed instructions Restructure document: Move discussions to extra sections Add FAQ, add analysis of used encoding space Add Pseudo-Ops, Macros, Algorithms Add Generalized Bit Permutations (shuffle)
2018-05-12	0.35	Replace shuffle with generalized zip (gzip) Add additional XBitfield ISA Extension Add figures and tables, Clean up document Extend discussion and evaluation chapters Add Verilog reference implementations Add fast C reference implementations

Date	Rev	Changes
2018-10-05	0.36	<hr/> XBitfield is now a proper extension proposal Add <code>bswaps.[hwd]</code> instructions Add <code>cmix</code> , <code>cmov</code> , <code>fsl</code> , <code>fsr</code> Rename <code>gzip</code> to <code>shfl/unshfl</code> Add <code>min</code> , <code>max</code> , <code>minu</code> , <code>maxu</code> Add <code>clri</code> , <code>maki</code> , <code>join</code> Add <code>cseln</code> , <code>cselz</code> , <code>mvnez</code> , <code>mveqz</code> Add <code>clmul</code> , <code>clmulh</code> , <code>bmatxor</code> , <code>bmator</code> , <code>bmatflip</code> Remove <code>bswaps.[hwd]</code> , <code>clri</code> , <code>maki</code> , <code>join</code> Remove <code>cseln</code> , <code>cselz</code> , <code>mvnez</code> , <code>mveqz</code> <hr/>
2019-06-10	0.90	<hr/> Add dedicated CRC instructions Add proposed opcode encodings Rename from XBitmanip to RISC-V Bitmanip Remove chapter on <code>bfxp[c]</code> instruction Refactor proposal into one big chapter Remove <code>c.brev</code> and <code>c.neg</code> instructions Add <code>fsri</code> , <code>pack</code> , <code>addiwu</code> , <code>slliu.w</code> Add <code>addwu</code> , <code>subwu</code> , <code>addu.w</code> , <code>subu.w</code> Rename <code>andc</code> to <code>andn</code> , Add <code>orn</code> and <code>xnor</code> Add <code>sbset[i]</code> , <code>sbclr[i]</code> , <code>sbinv[i]</code> , <code>sbext[i]</code> New naming scheme for <code>grevi</code> pseudo-ops Add <code>clmulr</code> instruction (reversed <code>clmul</code>) Jump to Rev 0.90 to indicate spec matureness <hr/>
2019-08-29	0.91	<hr/> Change encodings of <code>bmatxor</code> and <code>grev[i][w]</code> Add <code>gorc[i][w]</code> and <code>bfp[w]</code> instructions <hr/>
????-??-??	0.92	<hr/> — <hr/>

Bibliography

- [1] Apx/aux (pack/unpack) instructions on besm-6 mainframe computers. <http://www.mail.com.com/besm6/instset.shtml#pack>. Accessed: 2019-05-06.
- [2] Chess programming wiki, flipping mirroring and rotating. <https://chessprogramming.wikispaces.com/Flipping%20Mirroring%20and%20Rotating>. Accessed: 2017-05-05.
- [3] *MC88110 Second Generation RISC Microprocessor User's Manual*. Motorola Inc., 1991.
- [4] *Cray Assembly Language (CAL) for Cray X1 Systems Reference Manual*. Cray Inc., 2003. Version 1.1, S-2314-50.
- [5] *Cray XMT Principles of Operation*. Cray Inc., 2009. Version 1.3, S-2473-13.
- [6] *SPARC T3 Supplement to the UltraSPARC Architecture 2007 Specification*. Oracle, 2010.
- [7] *TMS320C64x/C64x+ DSP CPU and Instruction Set Reference Guide (Rev. J)*. Texas Instruments, 2010.
- [8] The Aggregate. The aggregate magic algorithms. <http://aggregate.org/MAGIC/>. Accessed: 2019-05-26.
- [9] Sean Eron Anderson. Bit twiddling hacks. <http://graphics.stanford.edu/~seander/bithacks.html>. Accessed: 2017-04-24.
- [10] Armin Biere. private communication, October 2018.
- [11] Vinodh Gopal, Erdinc Ozturk, Jim Guilford, Gil Wolrich, Wajdi Feghali, Martin Dixon, and Deniz Karakoyunlu. Fast crc computation for generic polynomials using pclmulqdq instruction. <https://www.intel.com/content/dam/www/public/us/en/documents/white-papers/fast-crc-computation-generic-polynomials-pclmulqdq-paper.pdf>, 2009. Intel White Paper, Accessed: 2018-10-23.
- [12] Y. Hilewitz, C. Lauradoux, and R. B. Lee. Bit matrix multiplication in commodity processors. In *2008 International Conference on Application-Specific Systems, Architectures and Processors*, pages 7–12, July 2008.
- [13] Yedidya Hilewitz and Ruby B. Lee. Fast bit compression and expansion with parallel extract and parallel deposit instructions. In *Proceedings of the IEEE 17th International Conference on Application-specific Systems, Architectures and Processors*, ASAP '06, pages 65–72, Washington, DC, USA, 2006. IEEE Computer Society.

- [14] James Hughes. Using carry-less multiplication (clmul) to implement erasure code. Patent US13866453, 2013.
- [15] Donald E. Knuth. *The Art of Computer Programming, Volume 4A*. Addison-Wesley, 2011.
- [16] Geoff Langdale and Daniel Lemire. Parsing gigabytes of JSON per second. *CoRR*, abs/1902.08318, 2019.
- [17] Daniel Lemire and Owen Kaser. Faster 64-bit universal hashing using carry-less multiplications. *CoRR*, abs/1503.03465, 2015.
- [18] George Marsaglia. Xorshift rngs. *Journal of Statistical Software, Articles*, 8(14):1–6, 2003.
- [19] Prashant Pandey, Michael A. Bender, and Rob Johnson. A fast x86 implementation of select. *CoRR*, abs/1706.00990, 2017.
- [20] Henry S. Warren. *Hacker’s Delight*. Addison-Wesley Professional, 2nd edition, 2012.
- [21] Wikipedia. Carry-less product. https://en.wikipedia.org/wiki/Carry-less_product. Accessed: 2018-10-05.
- [22] Wikipedia. Hamming weight. https://en.wikipedia.org/wiki/Hamming_weight. Accessed: 2017-04-24.
- [23] Wikipedia. Morton code (z-order curve, lebesgue curve). https://en.wikipedia.org/wiki/Z-order_curve. Accessed: 2018-10-12.
- [24] Clifford Wolf. Reference hardware implementations of bit extract/deposit instructions. <https://github.com/cliffordwolf/bextdep>. Accessed: 2017-04-30.
- [25] Clifford Wolf. Reference implementations of various crcs using carry-less multiply. <http://svn.clifford.at/handicraft/2018/clmulcrc/>. Accessed: 2018-11-06.
- [26] Clifford Wolf. A simple synthetic compiler benchmark for bit manipulation operations. <http://svn.clifford.at/handicraft/2017/bitcode/>. Accessed: 2017-04-30.