

RISC-V XBitmanip Extension
Document Version 0.34-draft

Editor: Clifford Wolf
Symbiotic GmbH
`clifford@symbioticeda.com`
April 9, 2018

Contributors to all versions of the spec in alphabetical order (please contact editors to suggest corrections): Steven Braeger, Rex McCrary, Po-wei Huang, and Clifford Wolf.

This document is released under a Creative Commons Attribution 4.0 International License.

Contents

1	Introduction	1
1.1	ISA Extension Proposal Design Criteria	1
1.2	B Extension Adoption Strategy	2
1.3	Next steps	2
2	RISC-V XBitmanip Extension	3
2.1	Count Leading Zeros (<code>clz</code>)	3
2.2	Count Bits Set (<code>pcnt</code>)	4
2.3	Generalized Reverse (<code>grev</code> , <code>grevi</code>)	4
2.4	Shift Ones (Left/Right) (<code>slo</code> , <code>sloi</code> , <code>sro</code> , <code>sroi</code>)	5
2.5	Rotate (Left/Right) (<code>rol</code> , <code>ror</code> , <code>rori</code>)	7
2.6	And-with-complement (<code>andc</code>)	8
2.7	Bit Extract/Deposit (<code>bext</code> , <code>bdep</code>)	9
2.8	Compressed instructions (<code>c.not</code> , <code>c.neg</code> , <code>c.clz</code> , <code>c.brev</code>)	10
3	Discussion and Evaluation	13
3.1	Frequently Asked Questions	13
3.2	Analysis of used encoding space	14
4	Pseudo-Ops, Macros, Algorithms	17
4.1	GREVI Pseudo-Ops	17
4.2	Bit Scanning Operations	18

4.3	MIX/MUX Operations	18
4.4	Butterfly Operations	19
5	Temporary encoding using NSE opcode space	21
6	Change History	25

Chapter 1

Introduction

This is the RISC-V XBitmanip Extension draft spec. Originally it was the B-Extension draft spec, but the work group got dissolved for bureaucratic reasons in November 2017.

It is currently an independently maintained document. We'd happily donate it to the RISC-V foundation as starting point for a new B-Extension work group, if there will be one.

1.1 ISA Extension Proposal Design Criteria

Any proposed changes to the ISA should be evaluated according to the following criteria.

- **Architecture Consistency:** Decisions must be consistent with RISC-V philosophy. ISA changes should deviate as little as possible from existing RISC-V standards (such as instruction encodings), and should not re-implement features that are already found in the base specification or other extensions.
- **Threshold Metric:** The proposal should provide a *significant* savings in terms of clocks or instructions. As a heuristic, any proposal should replace at least four instructions. An instruction that only replaces three may be considered, but only if the frequency of use is very high.
- **Data-Driven Value:** Usage in real world applications, and corresponding benchmarks showing a performance increase, will contribute to the score of a proposal. A proposal will not be accepted on the merits of its *theoretical* value alone, unless it is used in the real world.
- **Hardware Simplicity:** Though instructions saved is the primary benefit, proposals that dramatically increase the hardware complexity and area, or are difficult to implement, should be penalized and given extra scrutiny. The final proposals should only be made if a test implementation can be produced.
- **Compiler Support:** ISA changes that can be natively detected by the compiler, or are already used as intrinsics, will score higher than instructions which do not fit that criteria.

1.2 B Extension Adoption Strategy

The overall goal of this extension is pervasive adoption by minimizing potential barriers and ensuring the instructions can be mapped to the largest number of ops, either direct or pseudo, that are supported by the most popular processors and compilers. By adding generic instructions and taking advantage of the RISC-V base instructions that already operate on bits, the minimal set of instructions need to be added while at the same time enabling a rich set of operations.

The instructions cover the four major categories of bit manipulation: Count, Extract, Insert, Swap. The spec supports RV32, RV64, and RV128. “Clever” obscure and/or overly specific instructions are avoided in favor of more straight forward, fast, generic ones. Coordination with other emerging RISC-V ISA extensions groups is required to ensure our instruction sets are architecturally consistent.

1.3 Next steps

- Add support for this extension to processor cores and compilers so we can run quantitative evaluations on the instructions.
- Create assembler snippets for common operations that do not map 1:1 to any instruction in this spec, but can be implemented easily using clever combinations of the instructions. Add support for those snippets to compilers.

Chapter 2

RISC-V XBitmanip Extension

In the proposals provided in this section, the C code examples are for illustration purposes. They are not optimal implementations, but are intended to specify the desired functionality.

The sections on encodings are mere placeholders.

2.1 Count Leading Zeros (clz)

This operation counts the number of 0 bits before the first 1 bit (counting from the most significant bit) in the source register. This is related to the “integer logarithm”. It takes a single register as input and operates on the entire register. If the input is 0, the output is XLEN. If the input is ~ 0 , the output is 0.

```
uint_xlen_t clz(uint_xlen_t rs1)
{
    for (int count = 0; count < XLEN; count++)
        if ((rs1 << count) >> (XLEN-1))
            return count;
    return XLEN;
}
```

31	20 19	15 14	12 11	7 6	0
imm[11:0]		rs1	funct3	rd	opcode
12		5	3	5	7
???????????		src	CLZW	dest	OP-IMM-32
???????????		src	CLZ	dest	OP-IMM

One possible encoding for `clz` is as a standard I-type opcode somewhere in the brownfield surrounding the shift-immediate instructions.

2.2 Count Bits Set (pcnt)

The purpose of this instruction is to compute the number of 1 bits in a register. It takes a single register as input and operates on the entire register.

This operation counts the total number of set bits in the register.

```
uint_xlen_t pcnt(uint_xlen_t rs1)
{
    int count = 0;
    for(int index=0; index < XLEN; index++)
        count += (rs1 >> index) & 1;
    return count;
}
```

31	20 19	15 14	12 11	7 6	0
imm[11:0]	rs1	funct3	rd	opcode	
12	5	3	5	7	
??????????	src	PCNTW	dest	OP-IMM-32	
??????????	src	PCNT	dest	OP-IMM	

One possible encoding for `pcnt` is as a standard I-type opcode somewhere in the brownfield surrounding the shift-immediate instructions.

2.3 Generalized Reverse (grev, grevi)

The purpose of this instruction is to provide a single hardware instruction that can implement all of byte-order swap, bitwise reversal, short-order-swap, word-order-swap (RV64I), nibble-order swap, bitwise reversal in a byte, etc, all from a single hardware instruction. It takes in a single register value and an immediate that controls which function occurs, through controlling the levels in the recursive tree at which reversals occur.

This operation iteratively checks each bit `immed_i` from `i=0` to `XLEN-1`, in `XLEN` stages, and if the corresponding bit of the ‘function_select’ immediate is true for the current stage, swaps each adjacent pair of 2^i bits in the register.

`grevi` ‘butterfly’ implementation in C on various architectures

```
uint32_t grevi32(uint32_t rs1, int12_t immed)
{
    uint32_t x = rs1;
    uint5_t k = (uint12_t)immed;
    if (k & 1) x = ((x & 0x55555555) << 1) | ((x & 0xAAAAAAAA) >> 1);
    if (k & 2) x = ((x & 0x33333333) << 2) | ((x & 0xCCCCCCCC) >> 2);
}
```



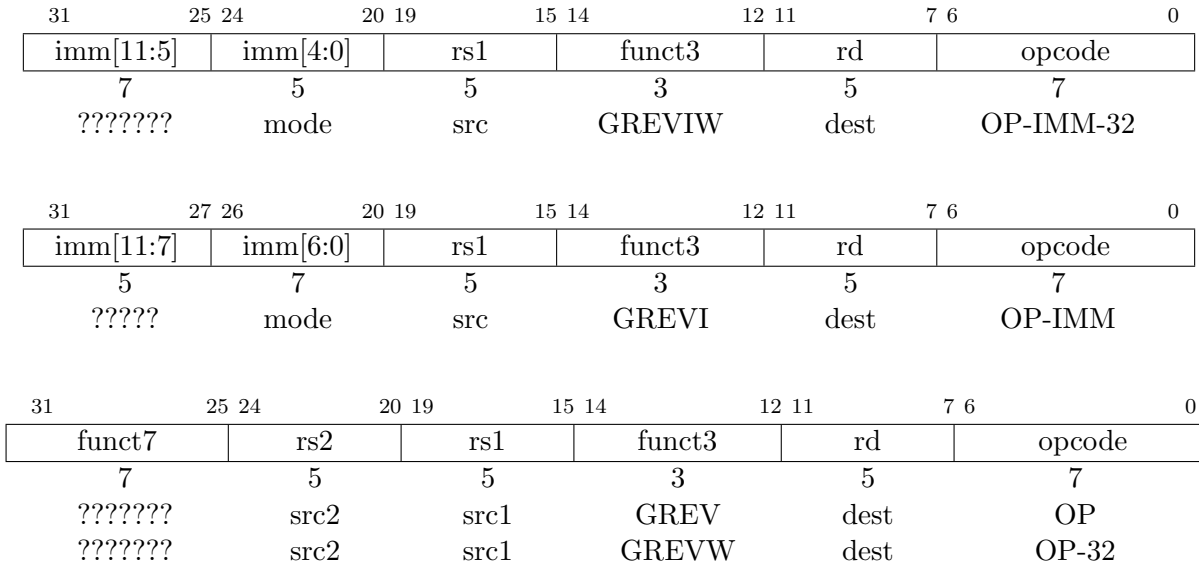
```

    if (k & 4) x = ((x & 0x0F0F0F0F) << 4) | ((x & 0xF0F0F0F0) >> 4);
    if (k & 8) x = ((x & 0x00FF00FF) << 8) | ((x & 0xFF00FF00) >> 8);
    if (k & 16) x = ((x & 0x0000FFFF) <<16) | ((x & 0xFFFF0000) >> 16);
    return x;
}

uint64_t grevi64(uint32_t rs1,int12_t immed)
{
    uint64_t x = rs1;
    uint6_t k = (uint12_t)imm;
    if (k & 1) x = ((x & 0x5555555555555555) << 1) | ((x & 0xAAAAAAAAAAAAAAAA) >> 1);
    if (k & 2) x = ((x & 0x3333333333333333) << 2) | ((x & 0xCCCCCCCCCCCCCCCC) >> 2);
    if (k & 4) x = ((x & 0x0F0F0F0F0F0F0F0F) << 4) | ((x & 0xF0F0F0F0F0F0F0F0) >> 4);
    if (k & 8) x = ((x & 0x00FF00FF00FF00FF) << 8) | ((x & 0xFF00FF00FF00FF00) >> 8);
    if (k & 16) x = ((x & 0x0000FFFF0000FFFF) <<16) | ((x & 0xFFFF0000FFFF0000) >> 16);
    if (k & 32) x = ((x & 0x00000000FFFFFF) <<32) | ((x & 0xFFFFFFFF00000000) >> 32);
    return x;
}

```

The above pattern should be intuitive to understand in order to extend this definition in an obvious manner for RV128+.



grev is encoded as standard R-type opcode and **grevi** is encoded as standard I-type opcode.

2.4 Shift Ones (Left/Right) (slo, sloi, sro, sroi)

These instructions are similar to shift-logical operations from the base spec, except instead of shifting in zeros, it shifts in ones. This can be used in mask creation or bit-field insertions, for example.

These instructions are exactly the same as the equivalent logical shift operations, except the shift shifts in ones values.

```
uint_xlen_t slo(uint_xlen_t rs1, uint_xlen_t rs2)
{
    int12_t amount=(rs2 & (XLEN-1));
    uint_xlen_t mask=((1 << amount)-1);
    return (rs1 << amount) | mask;
}
```

```
uint_xlen_t sloi(uint_xlen_t rs1, int12_t immed)
{
    int12_t amount=(immed & (XLEN-1));
    uint_xlen_t mask=((1 << amount)-1);
    return (rs1 << amount) | mask;
}
```

```
uint_xlen_t sro(uint_xlen_t rs1, uint_xlen_t rs2)
{
    int12_t amount=(rs2 & (XLEN-1));
    uint_xlen_t mask=~((1 << (XLEN-amount))-1);
    return (rs1 >> amount) | mask;
}
```

```
uint_xlen_t sroi(uint_xlen_t rs1, int12_t immed)
{
    int12_t amount=(immed & (XLEN-1));
    uint_xlen_t mask=~((1 << (XLEN-amount))-1);
    return (rs1 >> amount) | mask;
}
```

31	25 24	20 19	15 14	12 11	7 6	0
imm[11:5]	imm[4:0]	rs1	funct3	rd	opcode	
7	5	5	3	5	7	
10?????	shamt	src	SLOIW	dest	OP-IMM-32	
10?????	shamt	src	SROIW	dest	OP-IMM-32	

31	27 26	20 19	15 14	12 11	7 6	0
imm[11:7]	imm[6:0]	rs1	funct3	rd	opcode	
5	7	5	3	5	7	
10???	shamt	src	SLOI	dest	OP-IMM	
10???	shamt	src	SROI	dest	OP-IMM	

31	25 24	20 19	15 14	12 11	7 6	0
funct7	rs2	rs1	funct3	rd	opcode	
7	5	5	3	5	7	
10?????	src2	src1	SRO	dest	OP	
10?????	src2	src1	SLO	dest	OP	
10?????	src2	src1	SROW	dest	OP-32	
10?????	src2	src1	SLOW	dest	OP-32	

`s(l/r)o(i)` is encoded similarly to the logical shifts in the base spec. However, the spec of the entire family of instructions is changed so that the high bit of the instruction indicates the value to be inserted during a shift. This means that a `sloi` instruction can be encoded similarly to an `slli` instruction, but with a 1 in the highest bit of the encoded instruction. This encoding is backwards compatible with the definition for the shifts in the base spec, but allows for simple addition of a ones-insert.

When implementing this circuit, the only change in the ALU over a standard logical shift is that the value shifted in is not zero, but is a 1-bit register value that has been forwarded from the high bit of the instruction decode. This creates the desired behavior on both logical zero-shifts and logical ones-shifts.

2.5 Rotate (Left/Right) (`rol`, `ror`, `rori`)

These instructions are similar to shift-logical operations from the base spec, except they shift in the values from the opposite side of the register, in order. This is also called ‘circular shift’.

```
uint_xlen_t rol(uint_xlen_t rs1, uint_xlen_t rs2)
{
    int12_t amount=(rs2 & (XLEN-1));
    return (rs1 << amount) | (rs1 >> (XLEN-amount));
}

uint_xlen_t ror(uint_xlen_t rs1, uint_xlen_t rs2)
{
    int12_t amount=(rs2 & (XLEN-1));
    return (rs1 >> amount) | (rs1 << (XLEN-amount));
}

int_xlen_t rori(uint_xlen_t rs1, int12_t immmed)
{
    int12_t amount=(immmed & (XLEN-1));
    return (rs1 >> amount) | (rs1 << (XLEN-amount));
}
```

31	25 24	20 19	15 14	12 11	7 6	0
imm[11:5]	imm[4:0]	rs1	funct3	rd	opcode	
7	5	5	3	5	7	
11?????	shamt	src	RORIW	dest	OP-IMM-32	

31	27 26	20 19	15 14	12 11	7 6	0
imm[11:7]	imm[6:0]	rs1	funct3	rd	opcode	
5	7	5	3	5	7	
11???	shamt	src	RORI	dest	OP-IMM	

31	25 24	20 19	15 14	12 11	7 6	0
funct7	rs2	rs1	funct3	rd	opcode	
7	5	5	3	5	7	
11?????	src2	src1	RORW	dest	OP-32	
11?????	src2	src1	ROLW	dest	OP-32	
11?????	src2	src1	ROR	dest	OP	
11?????	src2	src1	ROL	dest	OP	

`ror(i),rol(i)` is implemented very similarly to the other shift instructions. One possible way to encode it is to re-use the way that bit 30 in the instruction encoding selects ‘arithmetic shift’ when bit 31 is zero (signalling a logical-zero shift). We can re-use this so that when bit 31 is set (signalling a logical-ones shift), if bit 31 is also set, then we are doing a rotate. The following table summarizes the behavior. The generalized reverse opcodes can be encoded using the bit pattern that would otherwise encode an “Arithmetic Left Shift” (which is an operation that does not exist).

Table 2.1: Rotate Encodings

Bit 31	Bit 30	Meaning
0	0	Logical Shift-Zeros
0	1	Arithmetic Shift
1	0	Logical Shift-Ones
1	1	Rotate

2.6 And-with-complement (andc)

31	25 24	20 19	15 14	12 11	7 6	0
funct7	rs2	rs1	funct3	rd	opcode	
7	5	5	3	5	7	
???????	src2	src1	ANDC	dest	OP	
???????	src2	src1	ANDCW	dest	OP-32	

This is an R-type instruction that implements the and-with-complement operation `x=(a & ~b)` with an assembly format of `andc ro,r1,r2`.

Additionally, the four positive bitwise operations can all be implemented with the same ALU area (e.g. as a single microcode instruction) using a scheme described in the implementation commentary, and fused versions can be implemented as well.

```
uint_xlen_t andc(uint_xlen_t rs1, uint_xlen_t rs2)
{
    return rs1 & ~rs2;
}
```

This instruction should be encoded similarly to the instruction in the base spec. The exact instruction encoding is to be decided, however.

2.7 Bit Extract/Deposit (bext, bdep)

This is an R-type instruction that implements the generic bit extract and bit deposit functions. Similar implementation can be referred to as gather/scatter or pack/unpack.

BEXT[W] rd,rs1,rs2 collects LSB justified bits to rd from rs1 using extract mask in rs2.

BDEP[W] rd,rs1,rs2 writes LSB justified bits from rs1 to rd using deposit mask in rs2.

```
uint_xlen_t bext(uint_xlen_t v, uint_xlen_t mask)
{
    uint32_t c = 0, m = 1;
    while (mask) {
        uint_xlen_t b = mask & -mask;
        if (v & b)
            c |= m;
        mask -= b;
        m <<= 1;
    }
    return c;
}

uint_xlen_t bdep32(uint_xlen_t v, uint_xlen_t mask)
{
    uint_xlen_t c = 0, m = 1;
    while (mask) {
        uint_xlen_t b = mask & -mask;
        if (v & m)
            c |= b;
        mask -= b;
        m <<= 1;
    }
    return c;
}
```

31	25 24	20 19	15 14	12 11	7 6	0
funct7	rs2	rs1	funct3	rd	opcode	
7	5	5	3	5	7	
???????	src2	src1	BEXT	dest	OP	
???????	src2	src1	BDEP	dest	OP	
???????	src2	src1	BEXTW	dest	OP-32	
???????	src2	src1	BDEPW	dest	OP-32	

This instruction should be encoded similarly to the instruction in the base spec. The exact instruction encoding is to be decided, however.

2.8 Compressed instructions (`c.not`, `c.neg`, `c.clz`, `c.brev`)

The RISC-V ISA has no dedicated instructions for bitwise inverse (`not`) and arithmetic inverse (`neg`). Instead `not` is implemented as `xori rd, rs, -1` and `neg` is implemented as `sub rd, x0, rs`.

In bitmanipulation code `not` and `neg` are very common operations. But there are no compressed encodings for those operations because there is no `c.xori` instructions and `c.sub` can not operate on `x0`.

Many bit manipulation operations that have dedicated opcodes in other ISAs must be constructed from smaller atoms in RISC-V XBitmanip code. But implementations might chose to implement them in a single micro-op using macro-op-fusion. For this it can be helpful when the fused sequences are short. `not` and `neg` are good candidates for macro-op-fusion, so it can be helpful to have compressed opcodes for them.

Likewise, `clz` and `brev` (an alias for `grevi rd, rs, -1`, i.e. bitwise reversal) are very common atoms for building common bit manipulation operations. So it is helpful to have compressed opcodes for them as well.

For example, the equivalent to the x86 instruction `BSR` (find index of highest set bit) can be implemented using the following 64 bit sequence that does not spill other registers (and thus is a prime candidate for macro-op-fusion):

```
bsr rd, rs:
    c.clz rd, rs
    c.neg rd
    addi rd, rd, 32
```

Likewise, the AMD TBM extension adds 10 x86 bit manipulation instructions. The equivalent RISC-V code sequences can utilize `c.not` in half of the cases.

The compressed instructions `c.not`, `c.neg`, `c.clz`, `c.brev` must be supported by all implementations that support the C extension and XBitmanip.

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0	
011			nzimm[9]				2				nzimm[4 6 8:7 5]			01		C.ADDI16SP (<i>RES</i> , <i>nzimm</i> = θ)
011			nzimm[17]				rd $\neq\{0, 2\}$				nzimm[16:12]			01		C.LUI (<i>RES</i> , <i>nzimm</i> = θ ; <i>HINT</i> , <i>rd</i> =0)
011			0		00		rs2'/rd'				0			01		C.NEG
011			0		01		rs1'/rd'				0			01		C.NOT
011			0		10		rs1'/rd'				0			01		C.CLZ
011			0		11		rs1'/rd'				0			01		C.BREV

This four instructions fit nicely in the reserved space in C.LUI/C.ADDI16SP. They only occupy 0.13% of the ≈ 15.6 bits wide RVC encoding space.

Chapter 3

Discussion and Evaluation

3.1 Frequently Asked Questions

Why `clz` instead of `ctz`? Why not `XYZ`?

The important thing is that we wanted to have a small ISA that only provides useful atoms for building most of the other "usual" bit manipulation instructions.

If we chose `clz` or `ctz` is not that important. But we want to pick only one and let the used build the other by combining the supported instruction with bitwise reverse (which is provided by `grevi`).

`grev[i]` seems to be overly complicated? Do we really need it?

The `grevi` instruction can be used to build a wide range of common bit permutation instructions. For example on RV32:

Table 3.1: Some GREVI Pseudo-Instructions

Mode	Meaning on RV32
4	swap the nibbles of each byte
7	reverse the bits in each byte
15	reverse the bits in each halfword
16	swap the two half words halfword
24	reverse the byte order
31	bitwise reverse

If `grevi` were removed from this spec we would need to add a few new instructions in its place for operations such as bitwise reverse or endianness conversion.

Do we really need all the *w opcodes for 32 bit ops on RV64?

I don't know. I think nobody does know at the moment. But they add very little complexity to the core. So the only question is if it is worth the encoding space. We need to run proper experiments with compilers that support those instructions. So they are in for now and if future evaluations show that they are not worth the encoding space then we can still throw them out.

Why only andc and not any other complement operators?

Early versions of this spec also included other *c operators. But experiments¹ have show that `andc` is much more common in bit manipulation code than any other operators. Especially because it is commonly used in `mix` and `mux` operations.

BEXT/BDEP look like really expensive operations. Do we really need them?

Yes, they are expensive, but not as expensive as one might expect. A single-cycle 32 bit BEXT+BDEP+GREV core can be implemented in less space than a single-cycle 16x16 bit multiplier with 32 bit output.²

3.2 Analysis of used encoding space

So how much encoding space is used by the XBitmanip extension?

Table 3.2: XBitmanip encoding space (\log_2 , i.e. in equivalent number of bits)

RV32		RV64		Instruction
2x	10	4x	10	CLZ, CLZW, PCNT, PCNTW
1x	15	3x	15	GREV, GREVW, GREVIW
1x	15	1x	16	GREVI
2x	15	6x	15	SLO, SRO, SLOW, SROW, SLOIW, SROIW
2x	15	2x	16	SLOI, SROI
2x	15	5x	15	ROR, ROL, RORW, ROLW, RORIW
1x	15	1x	16	RORI
3x	15	6x	15	ANDC, BEXT, BDEP, ANDCW, BEXTW, BDEPW
4x	4	4x	4	C.NEG, C.NOT, C.CLZ, C.BREV

In the compressed encoding space XBitmanip needs the equivalent of a 6 bit encoding space, or

¹<http://svn.clifford.at/handicraft/2017/bitcode/>

²<https://github.com/cliffordwolf/bextdep>

$\approx 0.13\%$ of the total ≈ 15.6 bits available. $(100/(2^{15.6-6}) \approx 0.13)$

On RV32, XBitmanip requires the equivalent of a ≈ 18.6 bit encoding space in the uncompressed encoding space. For comparison: A single standard I-type instruction (such as `ADDI` or `SLTIU`) requires a 22 bit encoding space. I.e. the entire RV32 XBitmanip extension needs less than one-eighth of the encoding space of the `SLTIU` instruction.

$$\frac{\log(2 \cdot 2^{10} + 12 \cdot 2^{15})}{\log(2)} \approx 18.6$$

On RV64, XBitmanip requires the equivalent of a ≈ 19.8 bit encoding space in the uncompressed encoding space. I.e. the entire RV64 XBitmanip extension needs less than one-quarter of the encoding space of the `SLTIU` instruction.

$$\frac{\log(4 \cdot 2^{10} + 20 \cdot 2^{15} + 4 \cdot 2^{16})}{\log(2)} \approx 19.8$$

Chapter 4

Pseudo-Ops, Macros, Algorithms

This chapter contains a collection of pseudo-ops, macros, and algorithms using the XBitmanip extension. For the sake of simplicity we assume RV32 for most examples in this chapter.

Most assembler routines in this chapter are written as if they were ABI functions, i.e. arguments are passed in a0, a1, ... and results are returned in a0. Registers t0, t1, ... are used for spilling.

4.1 GREVI Pseudo-Ops

Tables 4.1 and 4.2 list assembler pseudo-ops for commonly used `grevi` commands.

Pseudoinstruction	Base Instruction(s)	Meaning
<code>brev rd, rs</code>	<code>grevi rd, rs, 31</code>	bitwise reverse
<code>bswap rd, rs</code>	<code>grevi rd, rs, 24</code>	reverse the byte order
<code>bswap.h rd, rs</code>	<code>grevi rd, rs, 8</code>	reverse the byte order in each half word
<code>hswap rd, rs</code>	<code>grevi rd, rs, 16</code>	reverse the order of half words

Table 4.1: RV32 GREVI Pseudoinstructions.

Pseudoinstruction	Base Instruction(s)	Meaning
<code>brev rd, rs</code>	<code>grevi rd, rs, 63</code>	bitwise reverse
<code>bswap rd, rs</code>	<code>grevi rd, rs, 56</code>	reverse the byte order
<code>bswap.h rd, rs</code>	<code>grevi rd, rs, 8</code>	reverse the byte order in each half word
<code>bswap.w rd, rs</code>	<code>grevi rd, rs, 24</code>	reverse the byte order in each word
<code>hswap rd, rs</code>	<code>grevi rd, rs, 48</code>	reverse the order of half words
<code>hswap.w rd, rs</code>	<code>grevi rd, rs, 16</code>	reverse the order of half words in each word
<code>wswap rd, rs</code>	<code>grevi rd, rs, 32</code>	reverse the order of words

Table 4.2: RV64 GREVI Pseudoinstructions.

4.2 Bit Scanning Operations

Count Leading Zeros

This is trivially just the `clz` instruction:

```
clz a0, a0
```

Count Leading Ones

Simply invert before executing the `clz` instruction:

```
not a0, a0
clz a0, a0
```

Find index of highest bit set (aka `ilog2`)

```
clz a0, a0
neg a0, a0
addi a0, a0, 32
```

Find index of lowest bit set

```
brev a0, a0
clz a0, a0
```

4.3 MIX/MUX Operations

MIX Operation

A MIX operation selects bits from arguments `a1` and `a2` based on the bits in the control word `a0`.

```
and a1, a1, a0
andc a2, a2, a0
or a0, a1, a2
```

MUX Operation

A MUX operation selects word `a1` or `a2` based on if the control word `a0` is zero or nonzero, without branching. (The `snez` instruction is optional if `a0` is already either 0 or 1.)

```
snez a0, a0
neg a0, a0
and a1, a1, a0
andc a2, a2, a0
or a0, a1, a2
```

4.4 Butterfly Operations

The `grev`, `bext`, and `bdep` instructions are commonly implemented using butterfly circuits¹. (Butterfly circuits are a powerful tool for implementing bit permutations.) However, sometimes one wants to use butterfly operations directly in a program and it is not immediately obvious how to use GREVI and BDEP to perform butterfly operations.

One can easily implement a single butterfly stage using GREVI in just four instructions:

```
grevi t0, a0, N
and t0, t0, a1
andc t1, a0, a1
or a0, t0, t1
```

With `a0` containing the data input on entry and the result on exit. `N` is a power of two indicating which butterfly stage to implement, and `a1` contains an `XLEN` bitmask derived from the `XLEN/2` butterfly control word.

This bitmask can either be pre-computed or, if it needs to be generated on-the-fly, it can easily be calculated with the help of BDEP. For example for the first butterfly stage (using `a0` as input and output):

```
li t0, 0x55555555
bdep t0, a0, t0
slli a0, t0, 1
or a0, a0, t0
```

Likewise for the second butterfly stage:

```
li t0, 0x33333333
bdep t0, a0, t0
slli a0, t0, 2
or a0, a0, t0
```

For the third stage:

¹http://programming.sirrida.de/bit_perm.html#butterfly

```
li t0, 0x0f0f0f0f
bdep t0, a0, t0
slli a0, t0, 4
or a0, a0, t0
```

For the fourth stage:

```
li t0, 0x00ff00ff
bdep t0, a0, t0
slli a0, t0, 8
or a0, a0, t0
```

And for the fifth stage:

```
li t0, 0x0000ffff
and t0, a0, t0
slli a0, t0, 16
or a0, a0, t0
```


Chapter 5

Temporary encoding using NSE opcode space

This is currently a non-standard extension so we need instruction encodings using "custom" opcodes. The long-term goal for this spec is to become the official B extension. So the encoding in this chapter is a temporary solution for the meantime, allowing us to build compilers and processors to evaluate the extension.

The temporary encoding uses the funct3=000 space in opcode *custom-0* for the uncompressed instructions and the nzuimm[5]=1, nzuimm[4:0]=0 space in C.SLLI for the compressed RV32 instructions.

On RV64 there is no NSE space in the C instruction listing, therefore no encoding for compressed bit manipulation instructions is defined on RV64. (The availability of compressed bit manipulation instructions on RV32 should be sufficient to evaluate the utility of those instructions.)

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0	
000			nzuimm[5]					rs1/rd≠0							10	C.SLLI (<i>HINT, rd=0; RV32-NSE, nzuimm[5]=1</i>)
000			1					00							10	C.NEG
000			1					01							10	C.NOT
000			1					10							10	C.CLZ
000			1					11							10	C.BREV

31	27	26	25	24	20	19	15	14	12	11	7	6	0	
funct7				rs2		rs1		funct3		rd		opcode		R-type
imm[11:0]						rs1		funct3		rd		opcode		I-type

RV32 XBitmanip Instruction Set

0000000	00000	rs1	000	rd	0001011	CLZ
0001000	00000	rs1	000	rd	0001011	PCNT
0010000	rs2	rs1	000	rd	0001011	GREV
0010100	mode	rs1	000	rd	0001011	GREVI
0011000	rs2	rs1	000	rd	0001011	ANDC
0100000	rs2	rs1	000	rd	0001011	BEXT
0101000	rs2	rs1	000	rd	0001011	BDEP
1000100	shamt	rs1	000	rd	0001011	SLOI
1010100	shamt	rs1	000	rd	0001011	SROI
1110100	shamt	rs1	000	rd	0001011	RORI
1000000	rs2	rs1	000	rd	0001011	SLO
1010000	rs2	rs1	000	rd	0001011	SRO
1100000	rs2	rs1	000	rd	0001011	ROL
1110000	rs2	rs1	000	rd	0001011	ROR

31	27	26	25	24	20	19	15	14	12	11	7	6	0	
funct7				rs2		rs1		funct3		rd		opcode		R-type
imm[11:0]						rs1		funct3		rd		opcode		I-type

RV64 XBitmanip Instruction Set

0000000	00000	rs1	000	rd	0001011	CLZ
0000010	00000	rs1	000	rd	0001011	CLZW
0001000	00000	rs1	000	rd	0001011	PCNT
0001010	00000	rs1	000	rd	0001011	PCNTW
0010000	rs2	rs1	000	rd	0001011	GREV
0010010	rs2	rs1	000	rd	0001011	GREVW
001010	mode	rs1	000	rd	0001011	GREVI
0010110	mode	rs1	000	rd	0001011	GREVIW
0011000	rs2	rs1	000	rd	0001011	ANDC
0011010	rs2	rs1	000	rd	0001011	ANDCW
0100000	rs2	rs1	000	rd	0001011	BEXT
0100010	rs2	rs1	000	rd	0001011	BEXTW
0101000	rs2	rs1	000	rd	0001011	BDEP
0101010	rs2	rs1	000	rd	0001011	BDEPW
100010	shamt	rs1	000	rd	0001011	SLOI
1000110	shamt	rs1	000	rd	0001011	SLOIW
101010	shamt	rs1	000	rd	0001011	SROI
1010110	shamt	rs1	000	rd	0001011	SROIW
111010	shamt	rs1	000	rd	0001011	RORI
1110110	shamt	rs1	000	rd	0001011	RORIW
1000000	rs2	rs1	000	rd	0001011	SLO
1000010	rs2	rs1	000	rd	0001011	SLOW
1010000	rs2	rs1	000	rd	0001011	SRO
1010010	rs2	rs1	000	rd	0001011	SROW
1100000	rs2	rs1	000	rd	0001011	ROL
1100010	rs2	rs1	000	rd	0001011	ROLW
1110000	rs2	rs1	000	rd	0001011	ROR
1110010	rs2	rs1	000	rd	0001011	RORW

Chapter 6

Change History

Table 6.1: Summary of Changes

Date	Rev	Changes
2017-07-17	0.10	Initial Draft
2017-11-02	0.11	Removed roli, assembler can convert it to use a rori Removed bitwise subset and replaced with andc Doc source text same base for study and spec. Fixed typos
2017-11-30	0.32	Jump rev number to be on par with associated Study Moved pdep/pext into spec draft and called it scattergaher
2018-04-07	0.33	Move to github, throw out study, convert from .md to .tex Fixed typos and fixed some reference C implementations Rename bgat/bsca to bext/bdep Remove post-add immediate from clz Clean up encoding tables and code sections
????-??-??	0.34	Add GREV and compressed instructions Restructure document: Move discussions to extra sections Add FAQ, add analysis of used encoding space Add chapter on Pseudo-Ops, Macros, Algorithms Add chapter on temporary encoding using custom opcode space